ETH zürich

Distributed Computing Group (D-ITET)

Sybil Detection using GNNs

Master Thesis Presentation

Stuart Heeb

September 12th, 2024

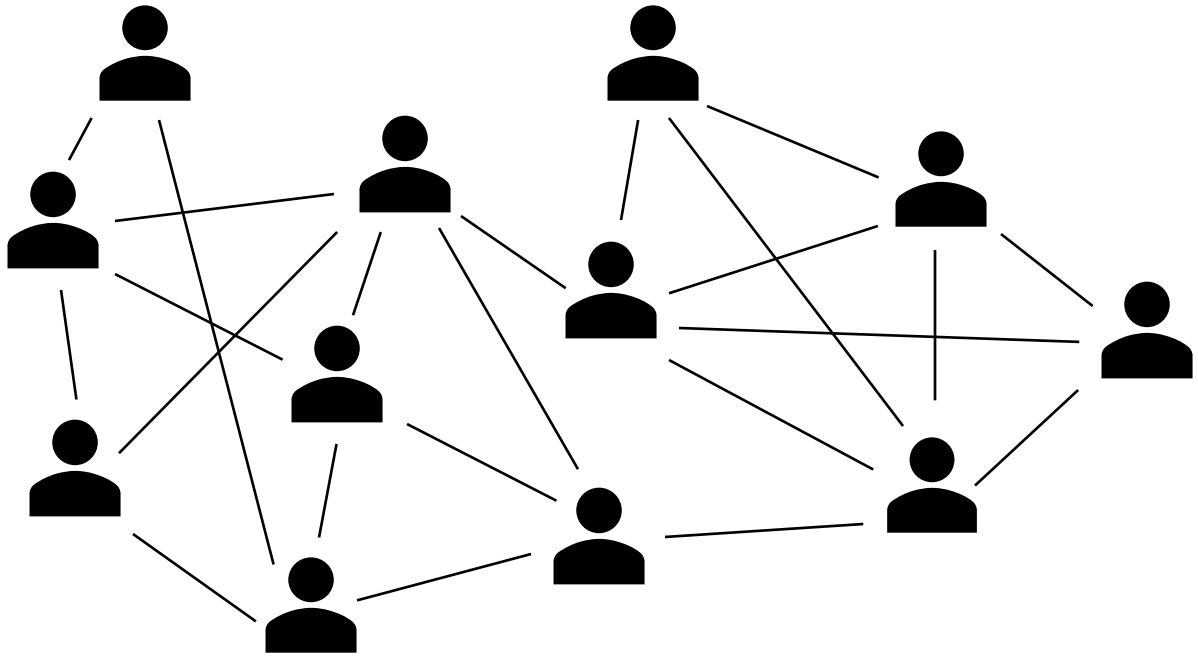

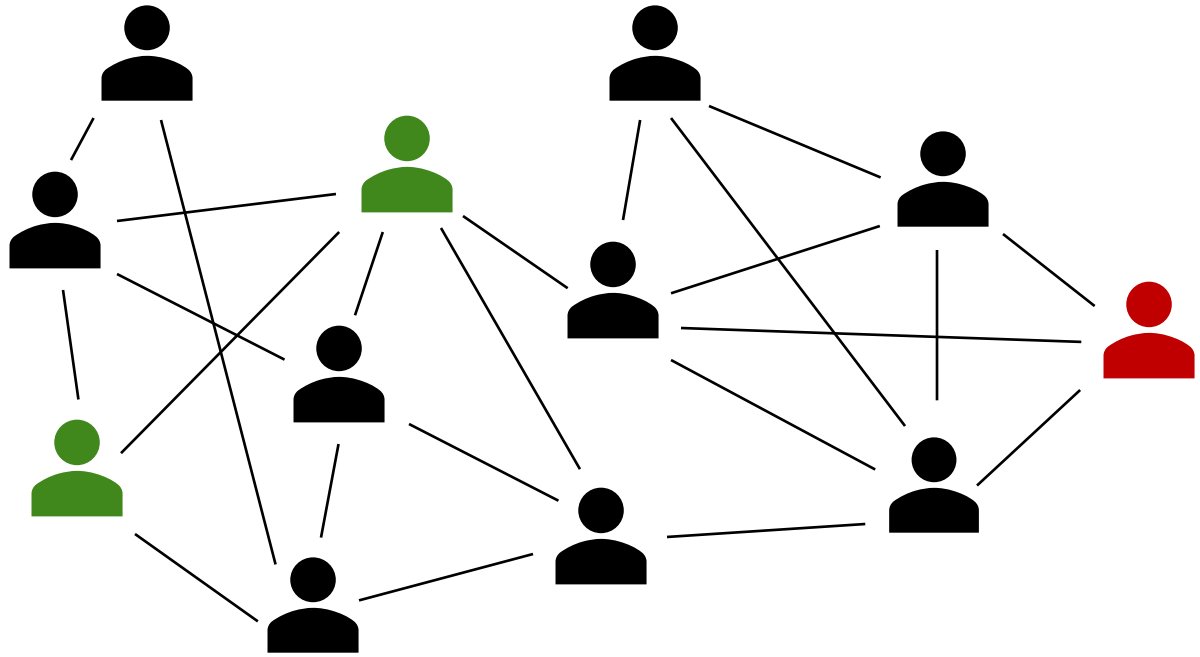

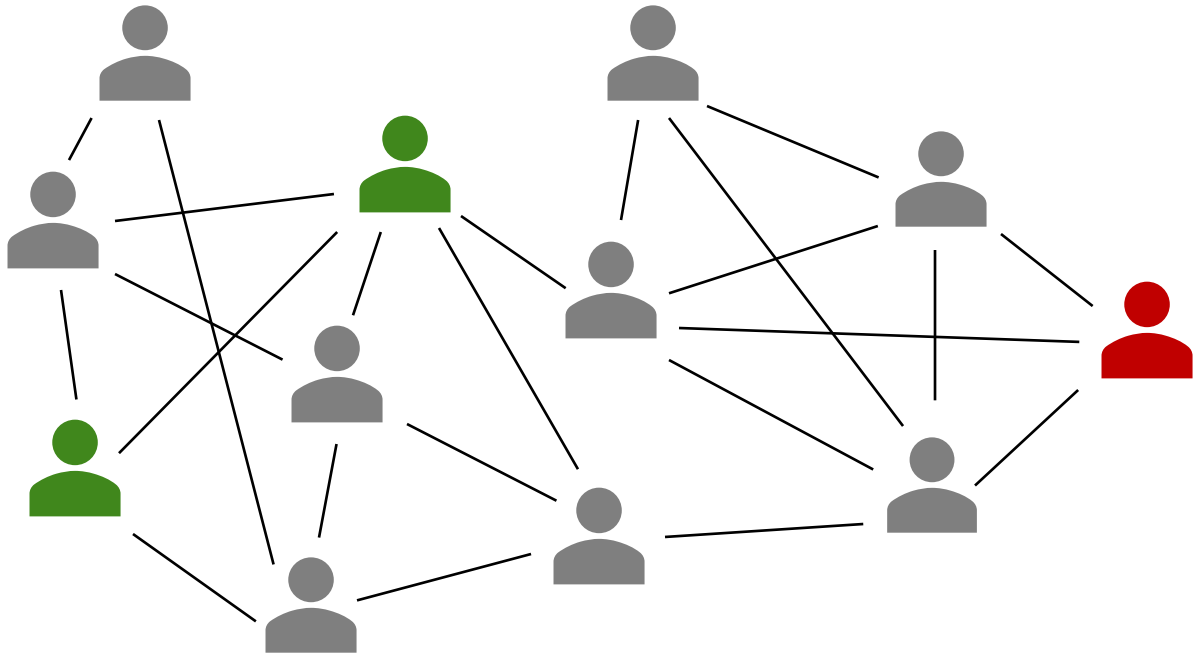

Known honest nodes

Known sybil nodes

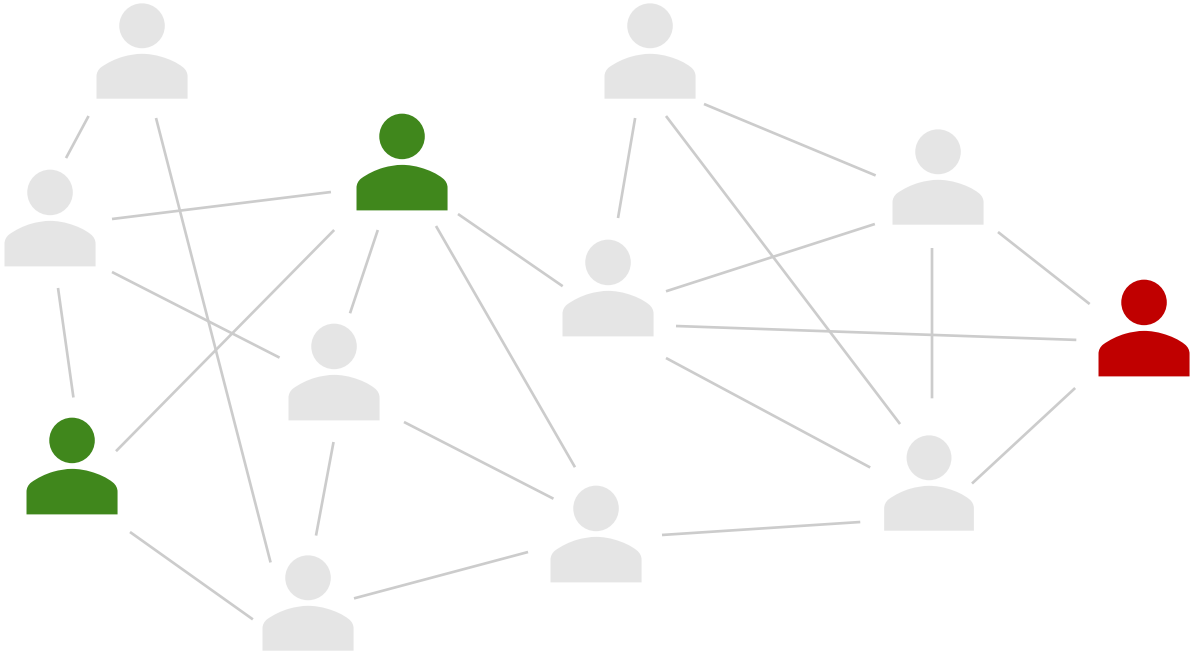

… Sybil Detection Algorithm

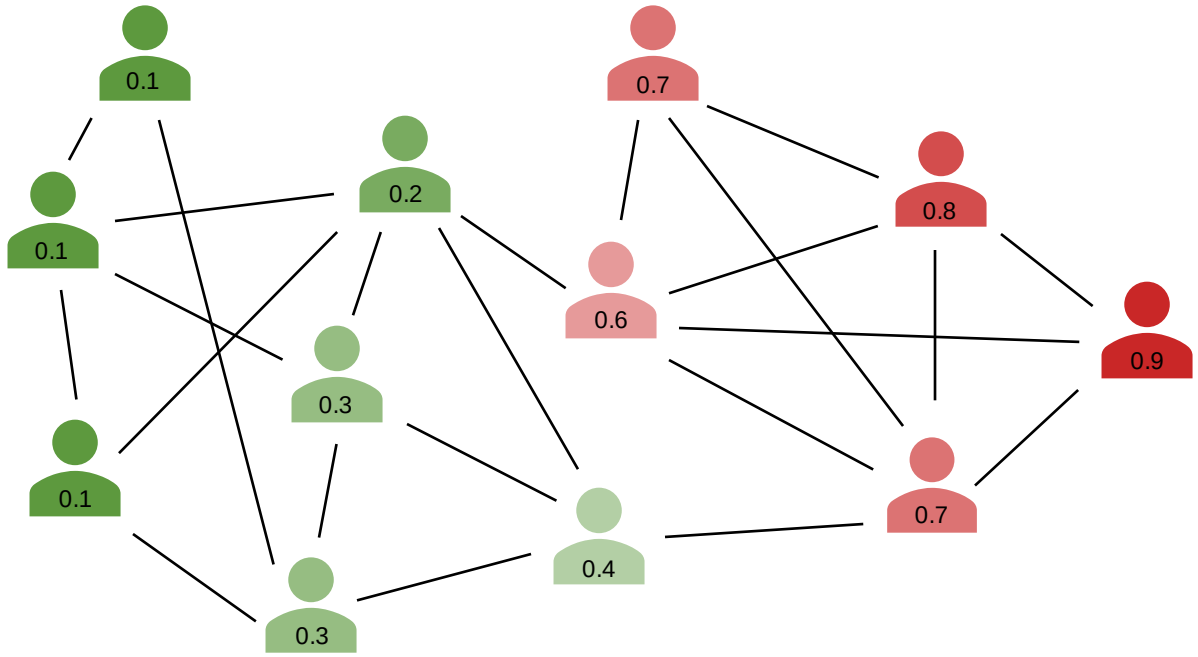

$$P(x_u = 1 \mid \mathbf{x}_V)$$

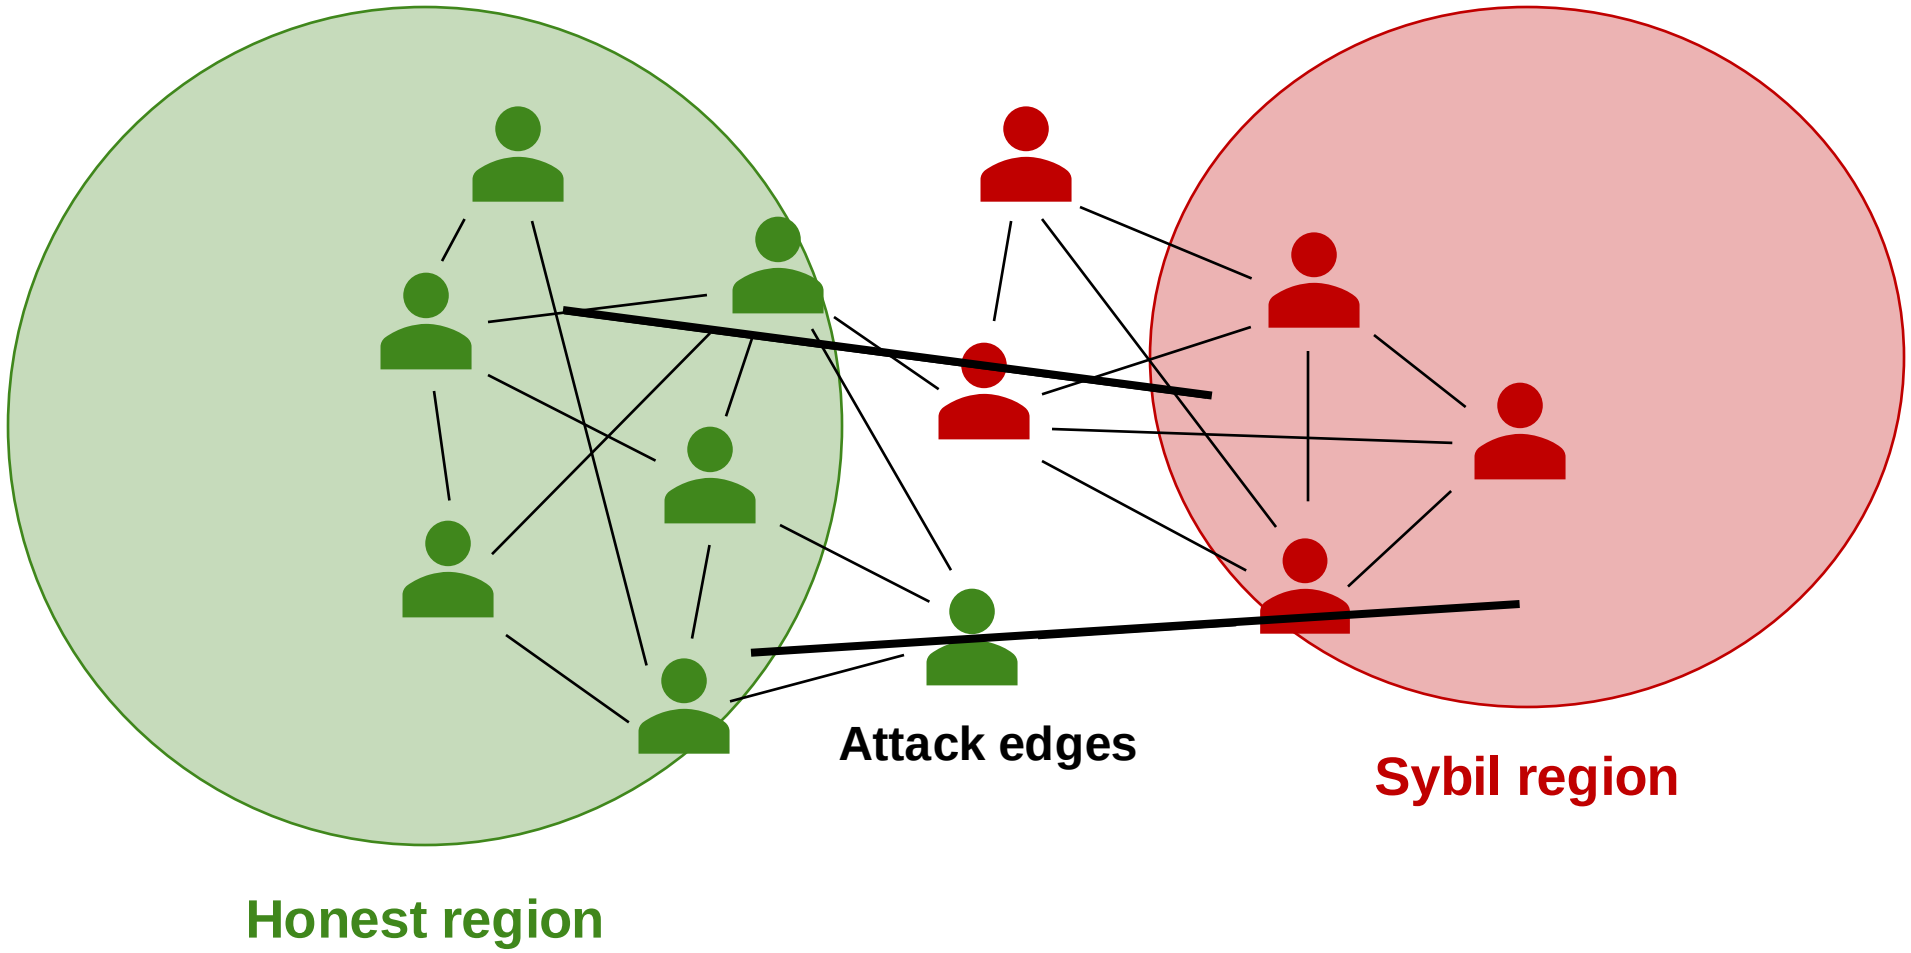
Attack edges
Sybil region
Honest region

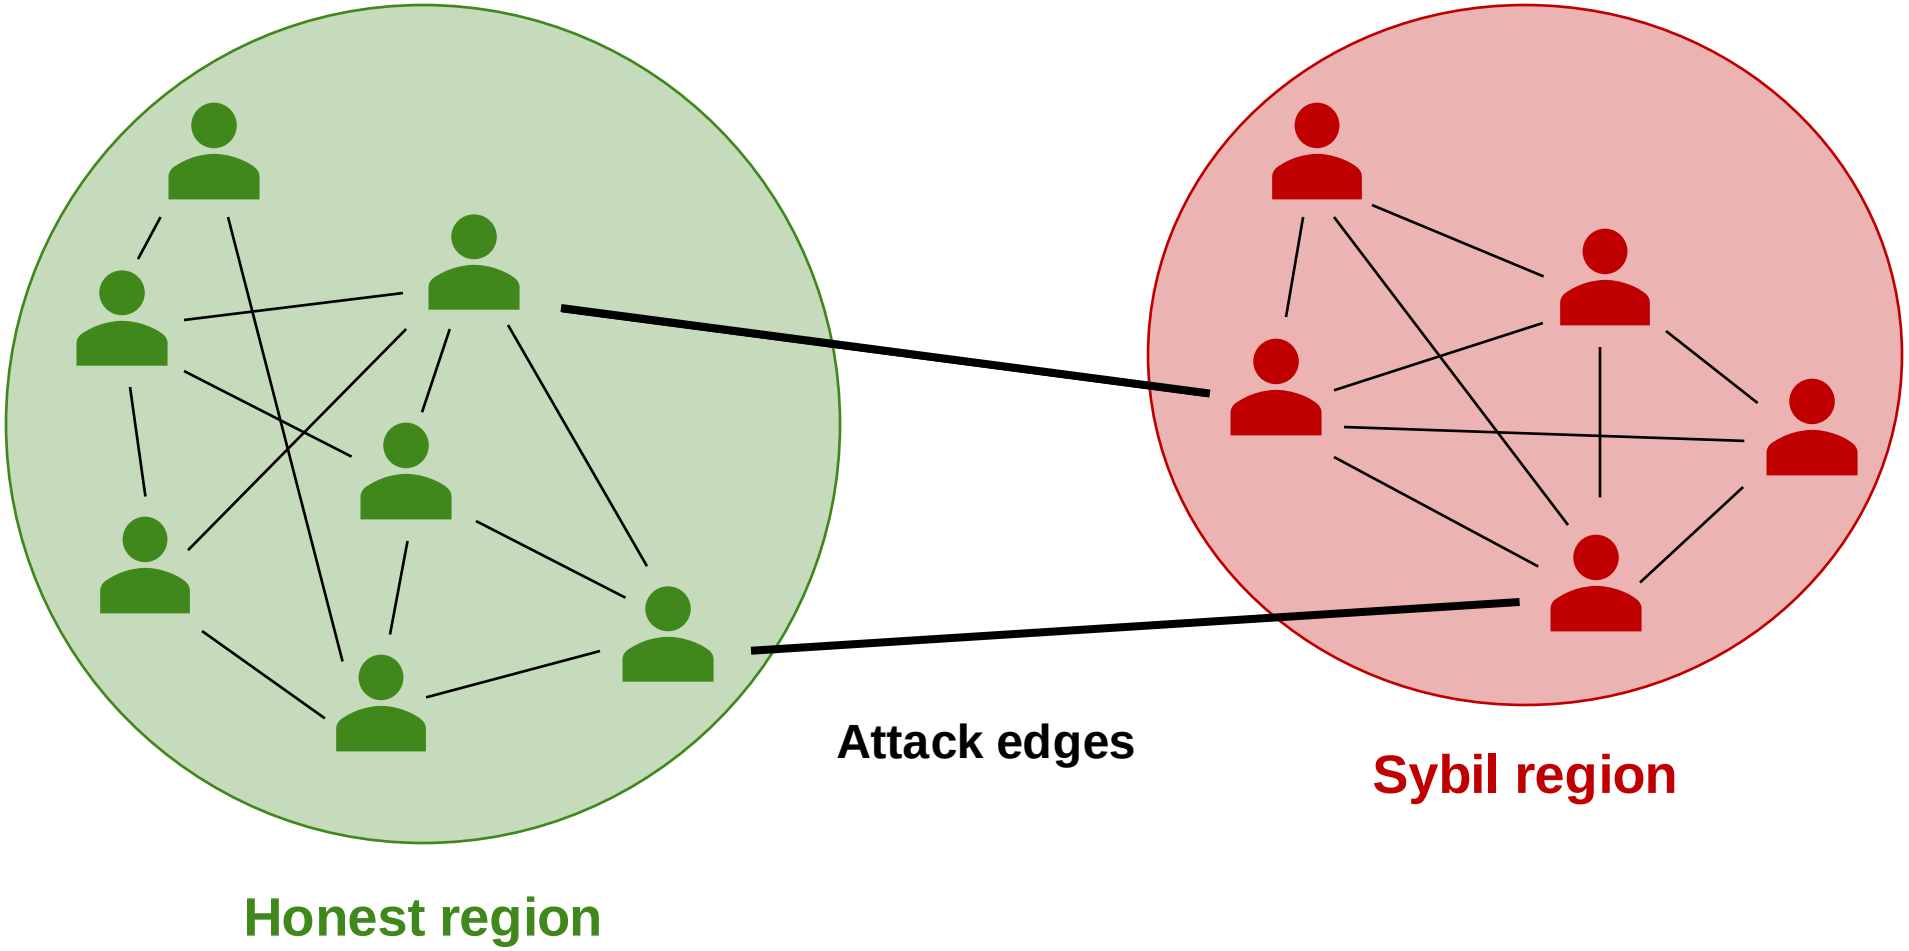
Attack edges
Sybil region
Honest region

Why are Sybils a Problem?

Spread misinformation

Degrade trust

Influence public opinion

https://www.cnet.com/pictures/ai-or-not-ai-can-you-spot-the-real-photos/

Modern Sybils

AI-generated

Actual photograph

Which picture is real?

As in, an image of the real world taken by a physical camera?

Structure-based Sybil Detection

- Content becoming harder to (even manually) distinguish
- Other advantages
  - Privacy concerns
  - Enhanced generalizability

Background & Related Work

SybilRank

2012

- Random Walks (RW)
- Uses only honest labels
- Computationally efficient

SybilBelief

2014

- Loopy Belief Propagation (LBP)
- Uses both honest and sybil labels

SybilSCAR

2019

- Local rule-based propagation
- Uses both honest and sybil labels
- Combines benefits of RW and LBP

Why use GNNs?

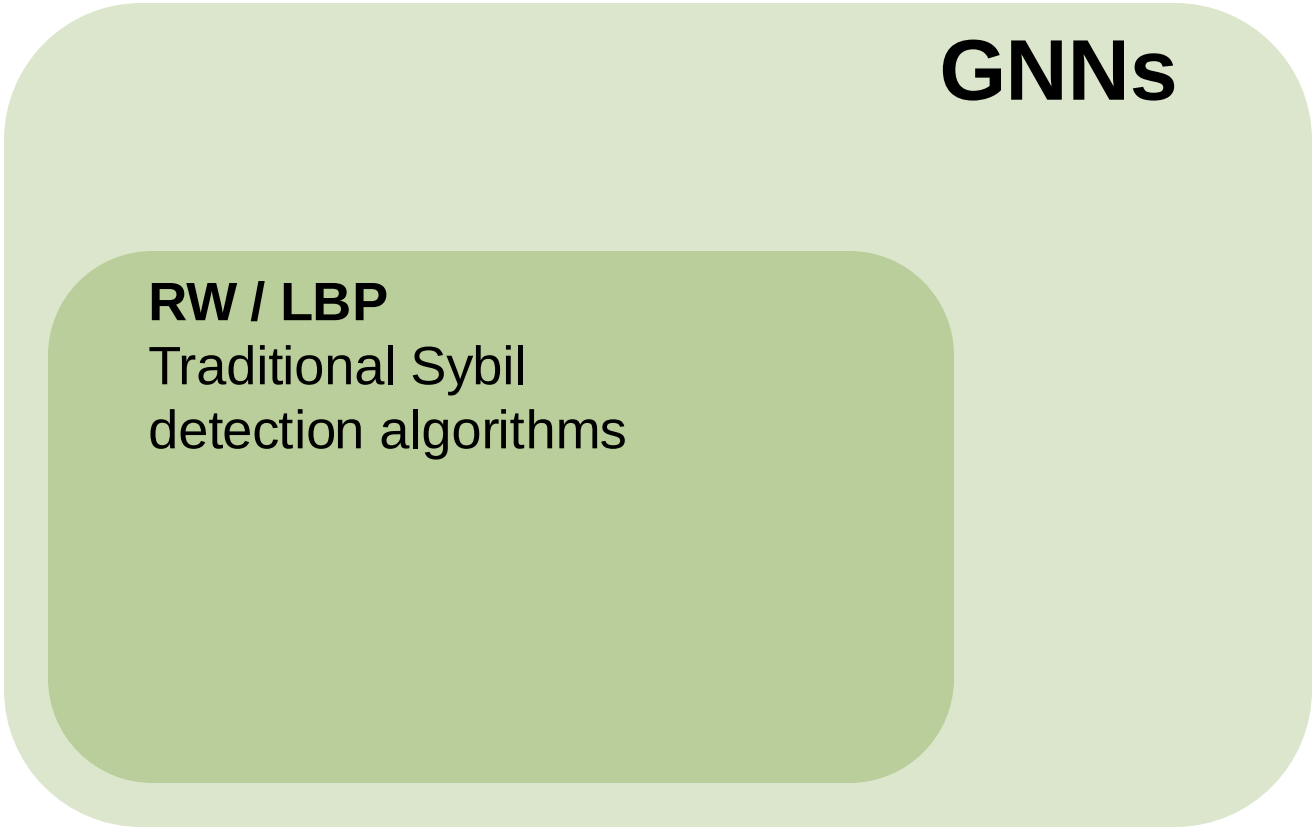

GNNs as a «one size fits all»?

GNNs for Sybil Detection

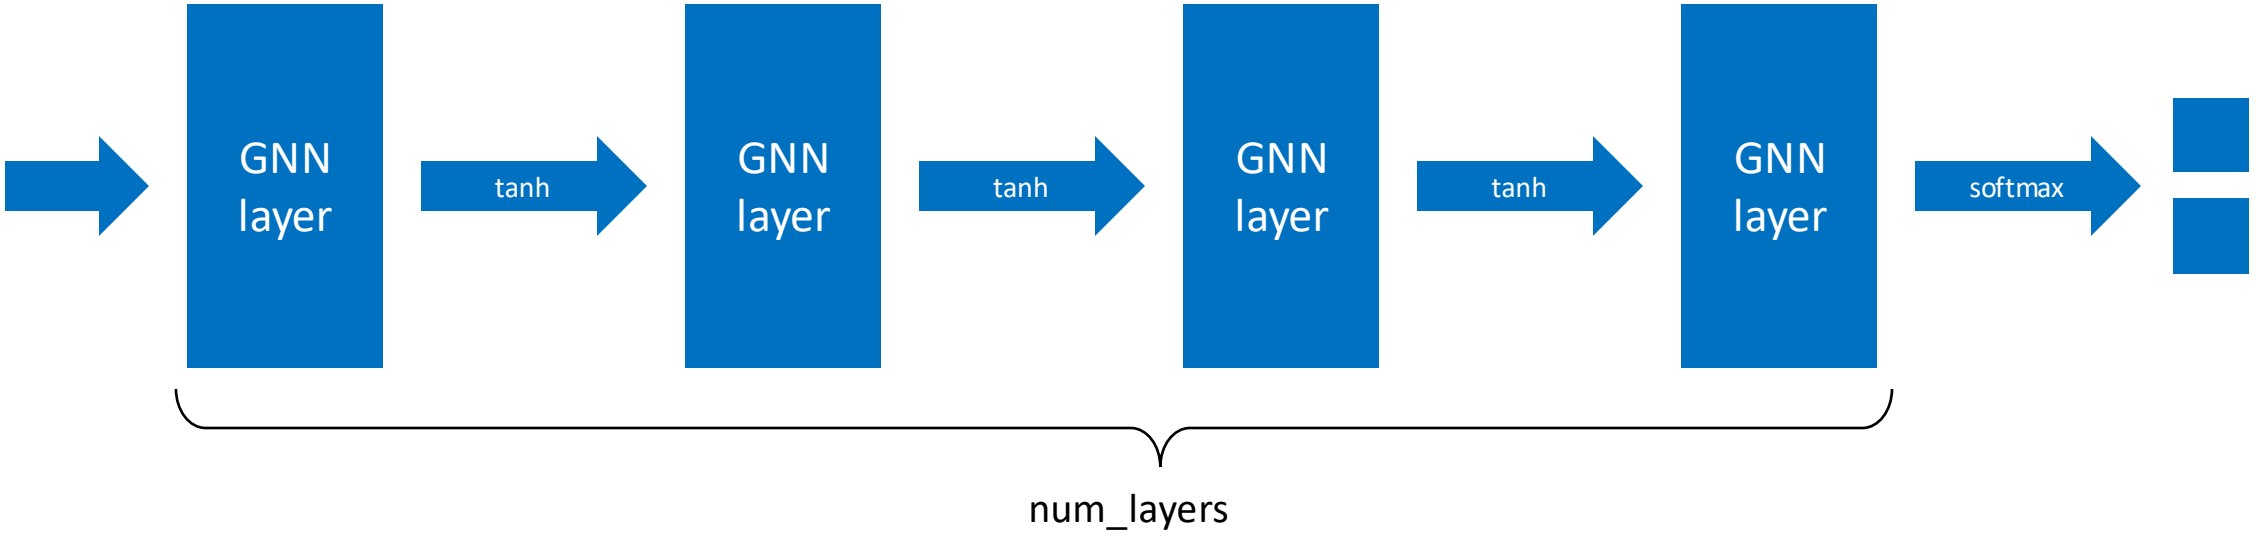

SybilGCN SybilRGCN SybilGAT

GNNs for Sybil Detection
Approaches

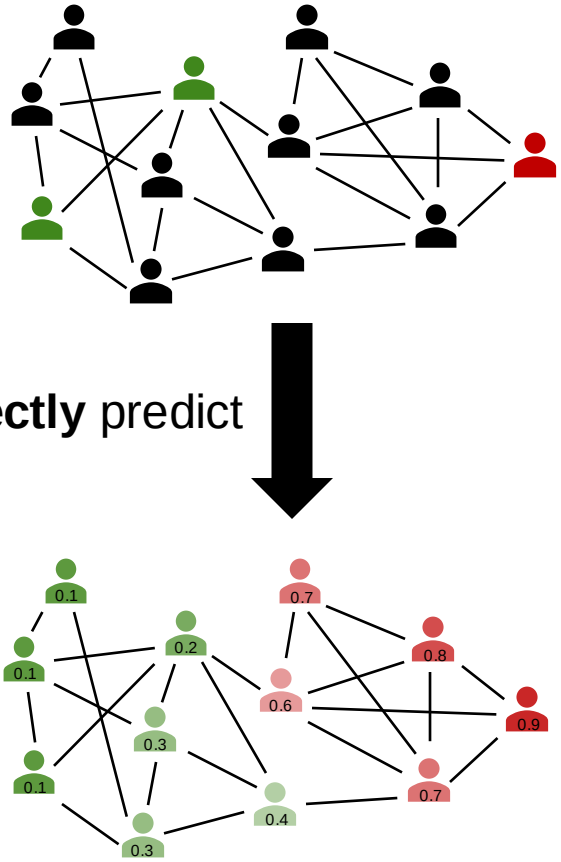

Without pre-training

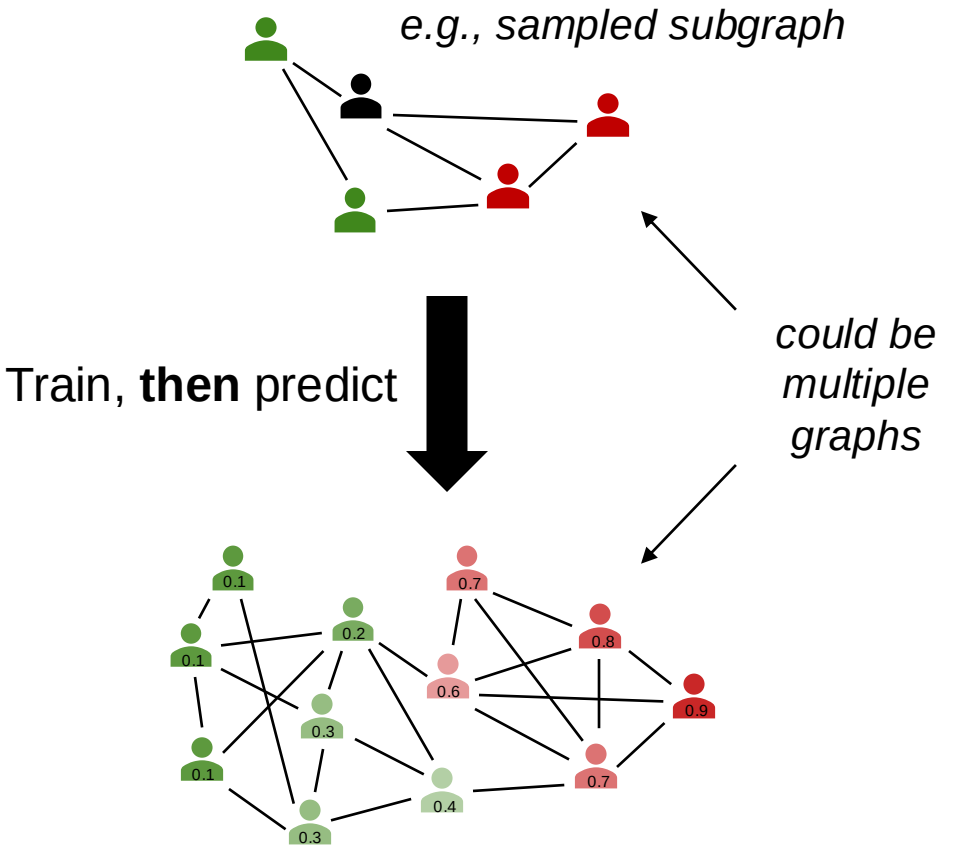

With pre-training

Data

270K nodes
6.8M edges
labeled

4K nodes
88K edges
unlabeled

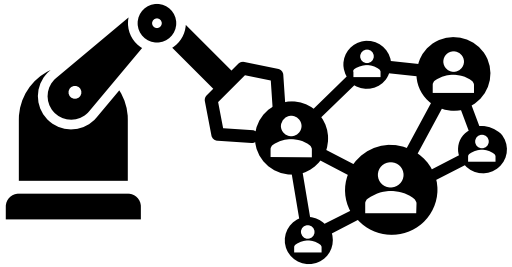

Synthetic
networks

https://www.vecteezy.com/free-vector/facebook-logo-black
https://www.freepik.com/free-photos-vectors/x-logo

Results: Performance Comparison

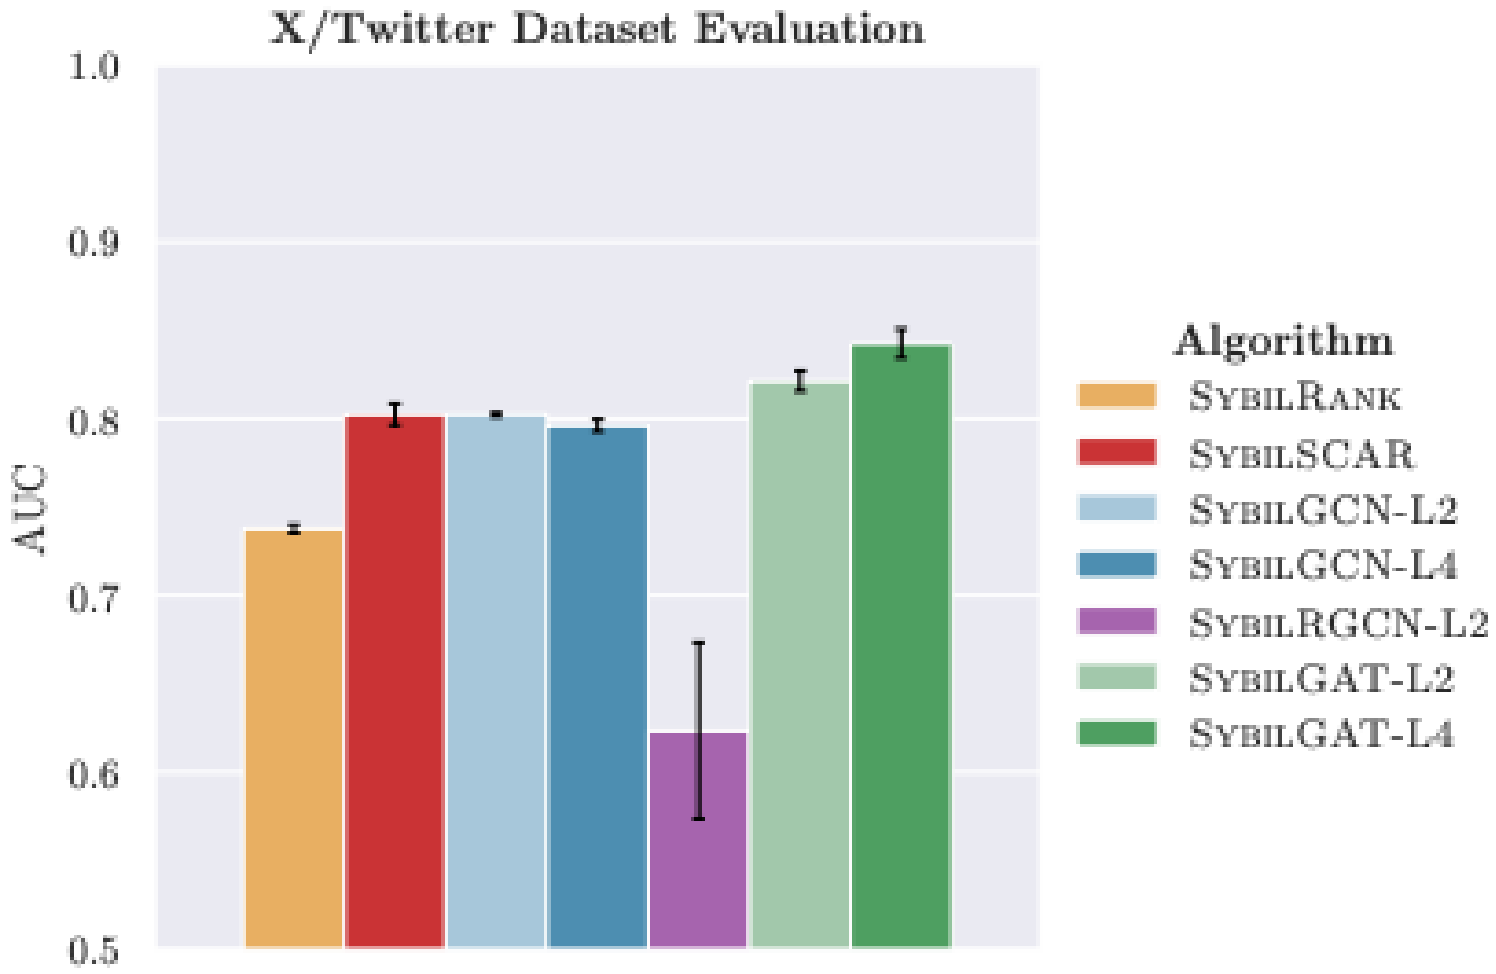

Results: Social Networks Synthetization

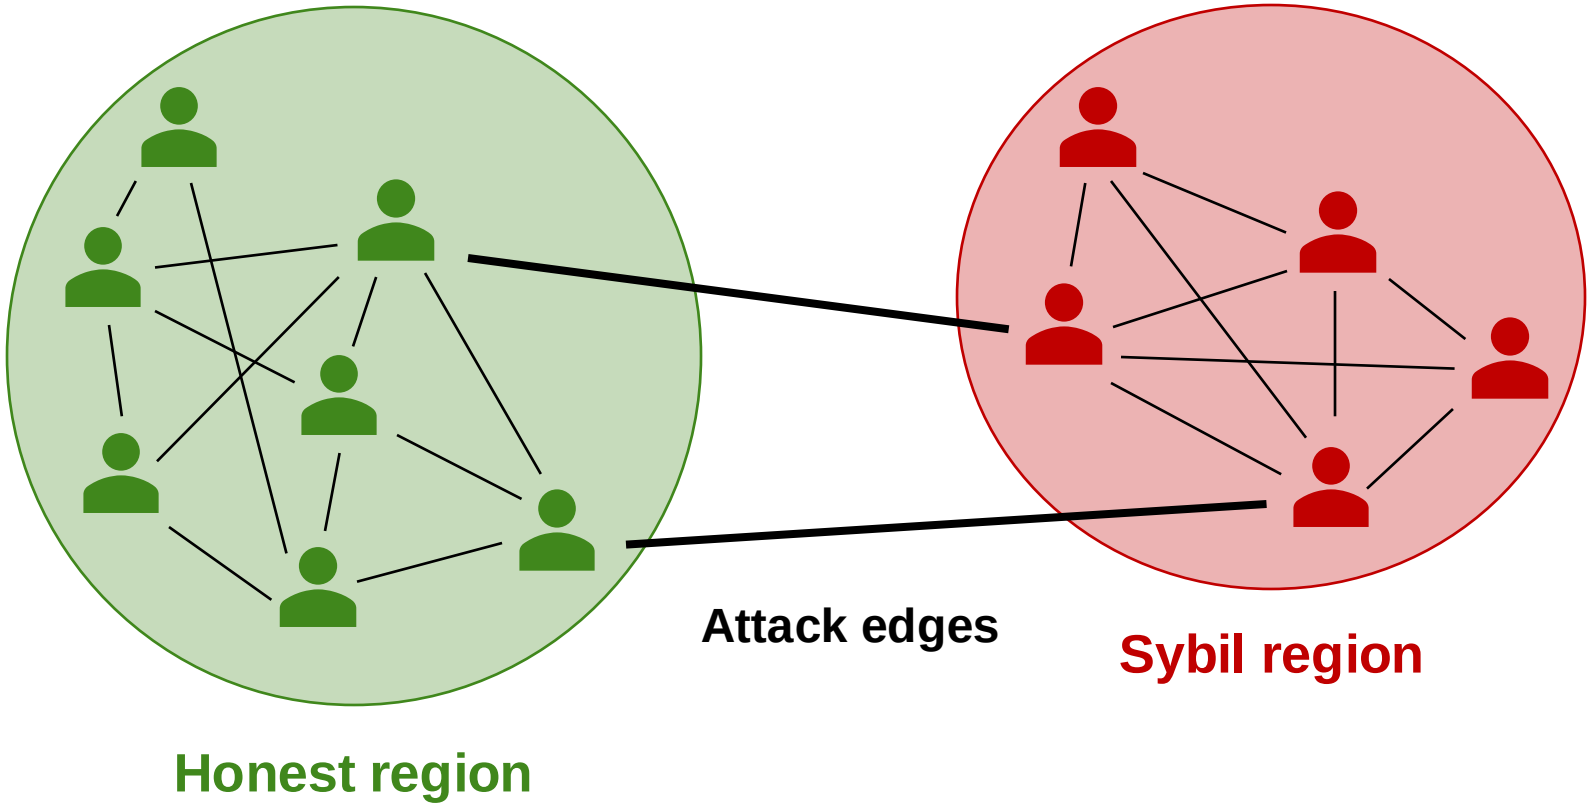

How to synthesize attack edges?

Results: Robustness
Label Noise Level

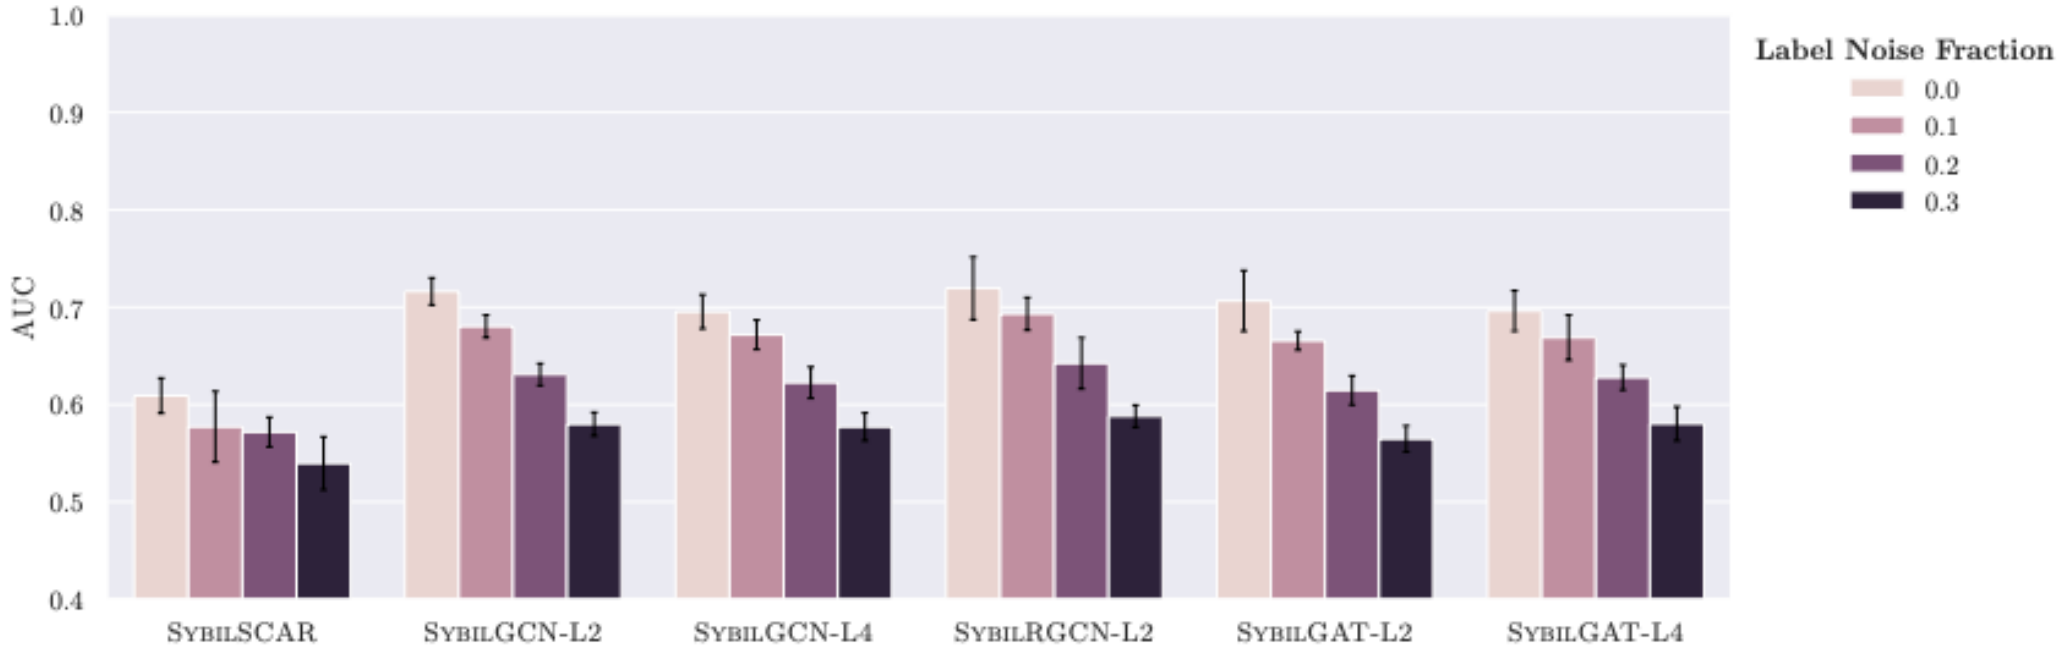

Results: Robustness
Label Noise Level

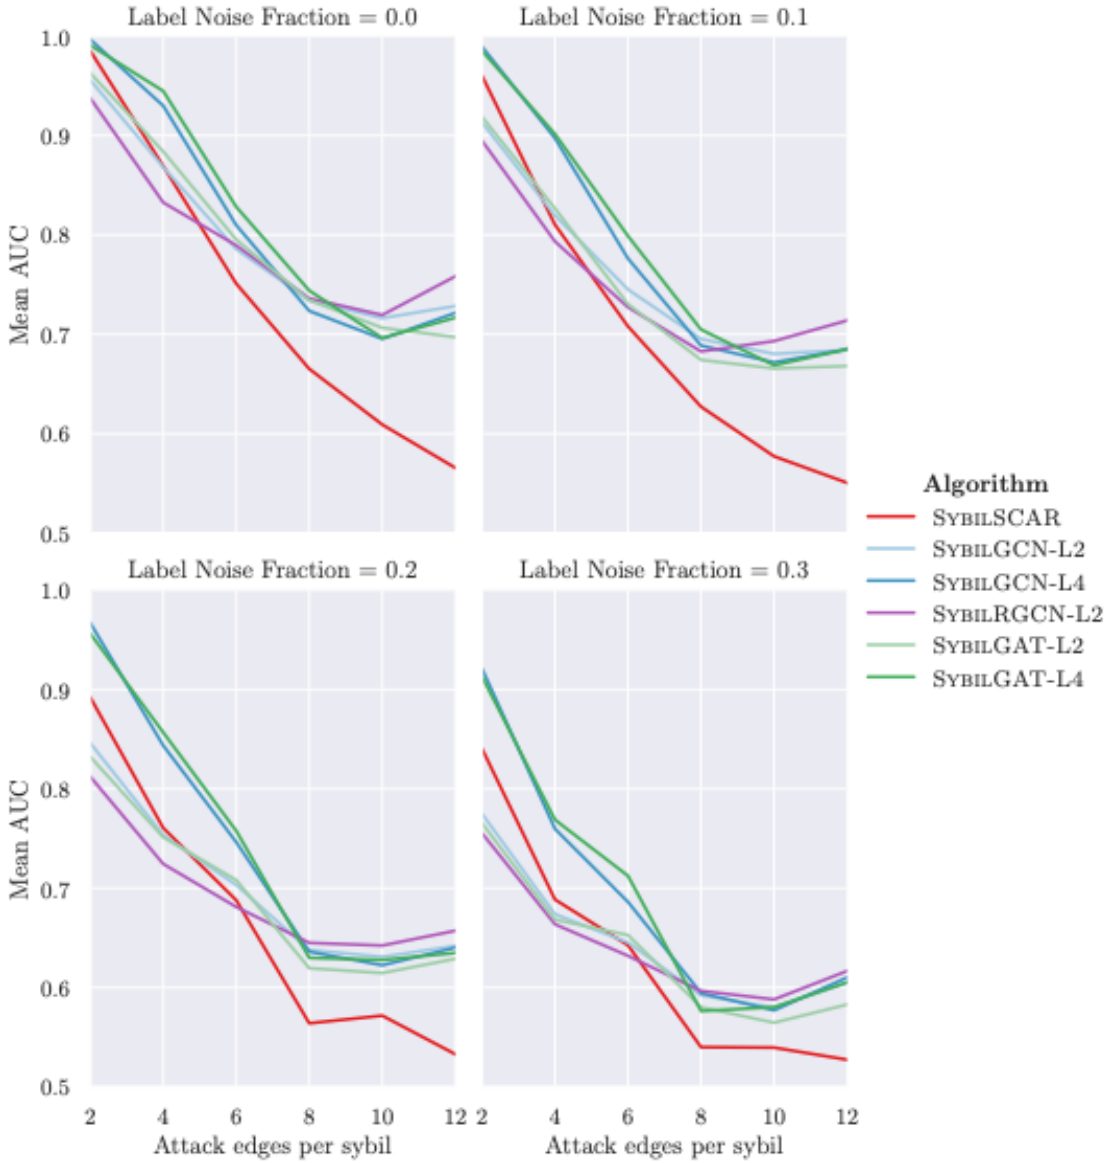

Adversarial Attack

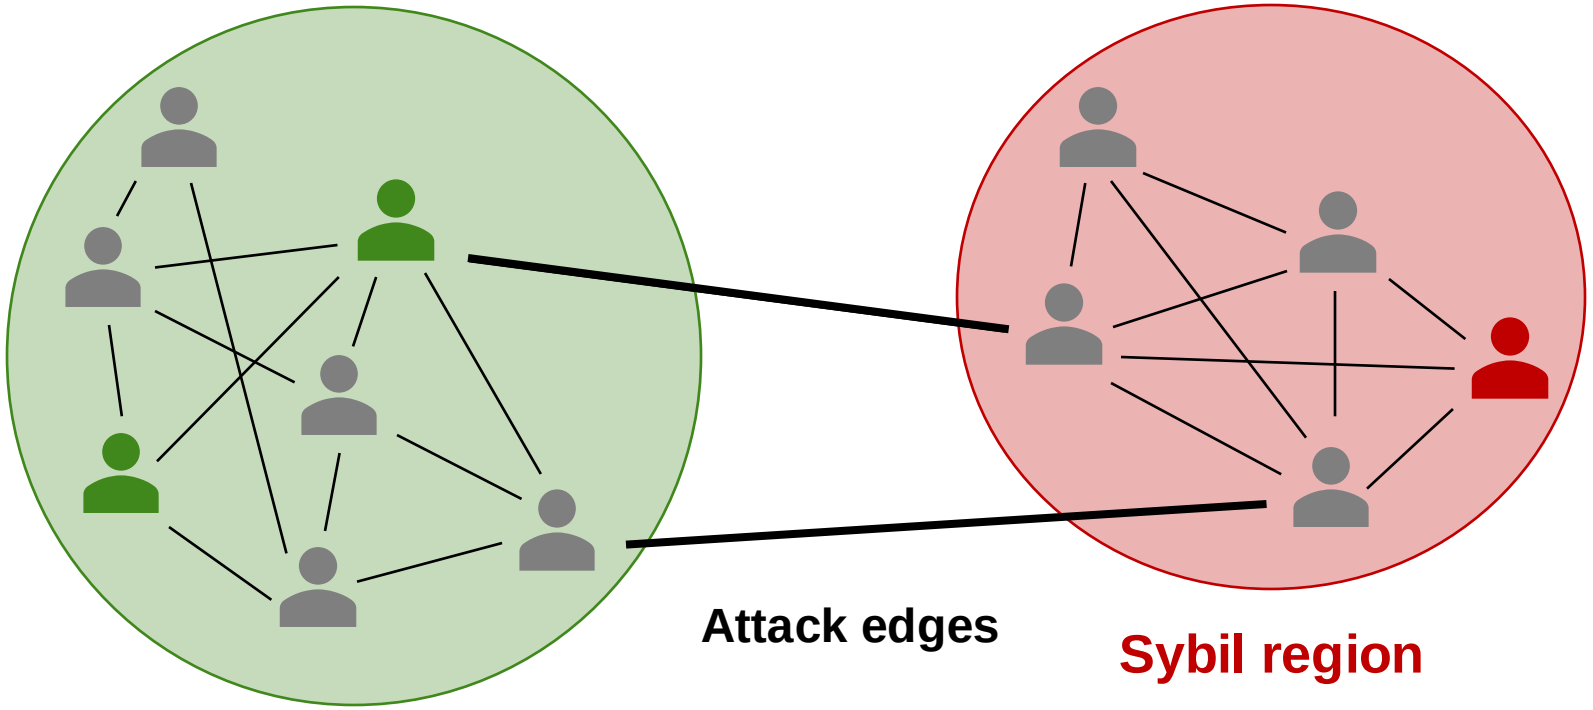

How to place attack edges in a targeted fashion?

$$p_t \qquad \boldsymbol{p}$$

Adversarial Attack

Targeted attack edge

With prob. p_t

$\cap$

$\text{Uniform}(N_k(\ \mathbf{H}_{\text{train}}\ ))$

With prob. p_k

$\cap$

$\text{Uniform}(\ \mathbf{S}\ )$

$$\boldsymbol{p} = [0.25, 0.25, 0.5] \in [0, 1]^K$$

direct hit neighbor 2-hop neighbor

Results: Adversarial Attacks

Increasing expected number of targeted attack edges

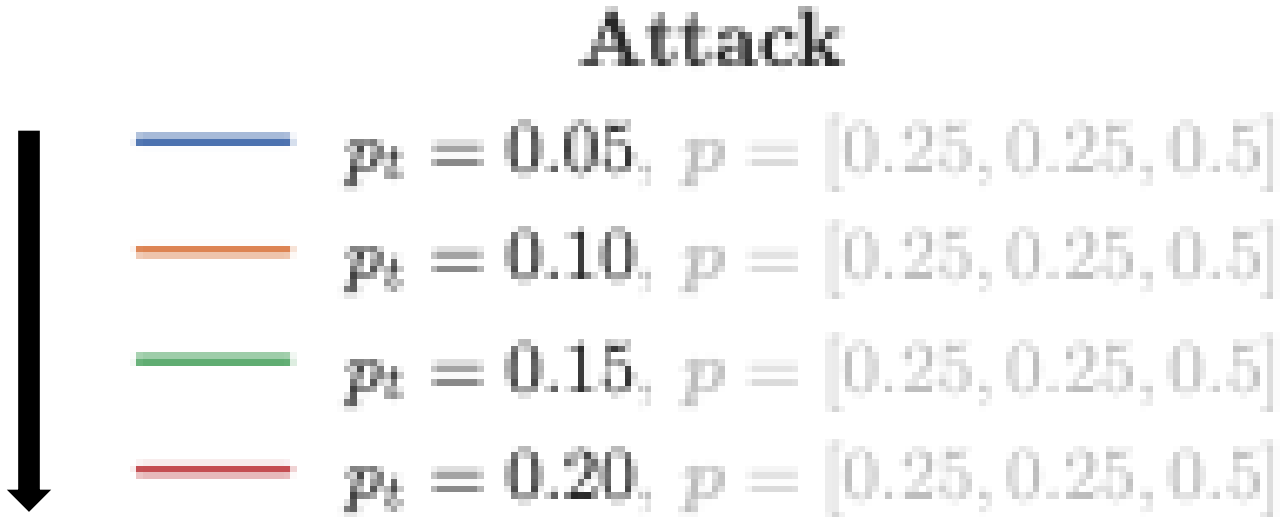

Results: Adversarial Attacks

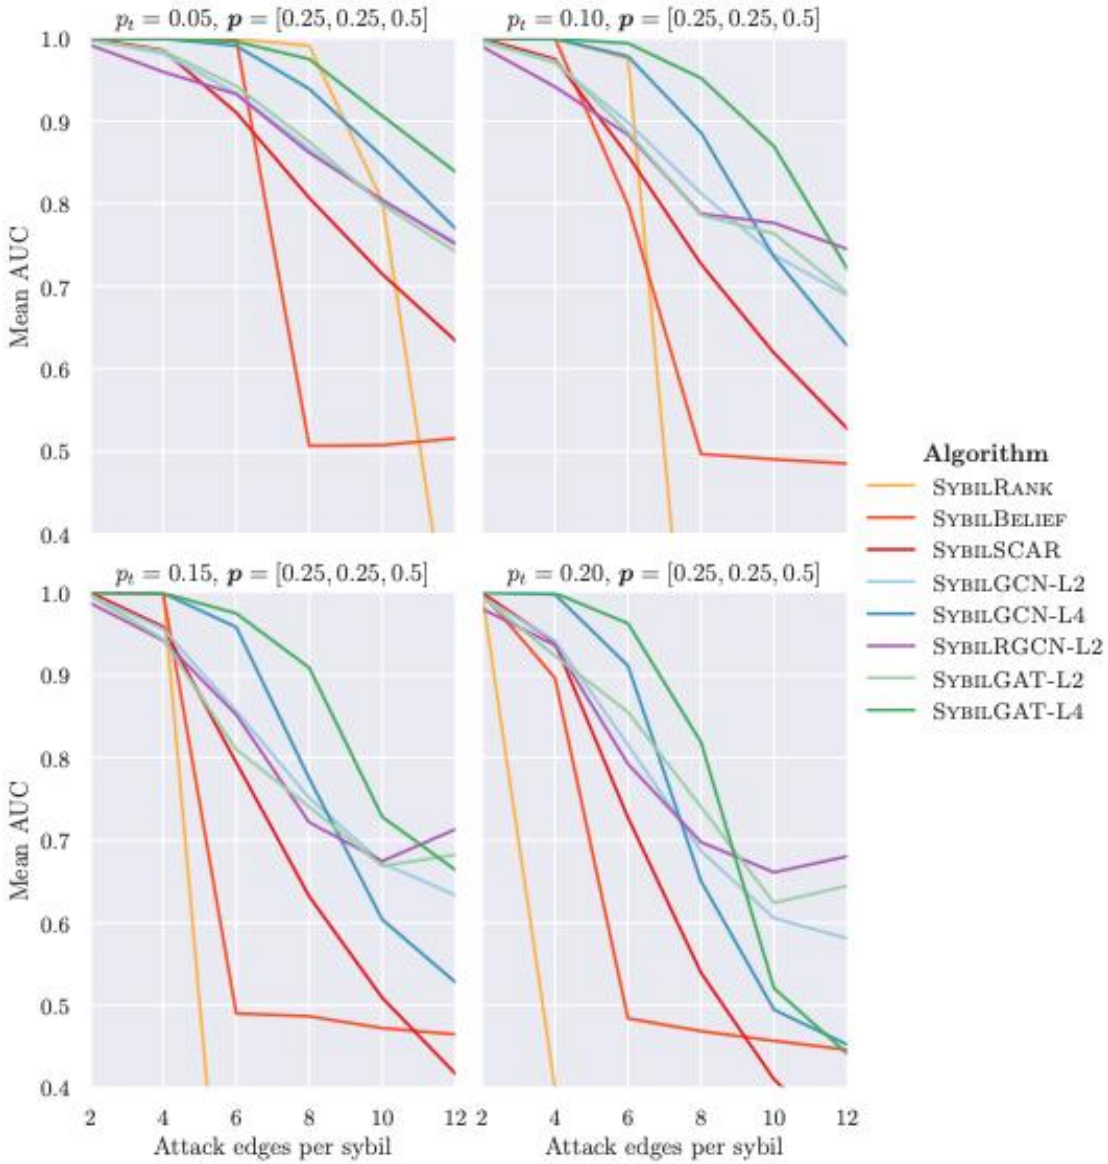

Conclusion

- SybilGCN and SybilGAT algorithms outperform the baselines in almost all scenarios, including real-world X/Twitter dataset
- SybilRGCN excels when attack complexity is high
- Synthesized social networks

Limitations & Future Work

- More advanced GNN architecture somewhat unexplored
- Robustness to lack of training data and label noise
- More data
- Theoretical analysis

ETH zürich

Thank you for listening!

D ITET D INFK

Stuart Heeb

stuart.heeb@inf.ethz.ch

Backup Slides


trust

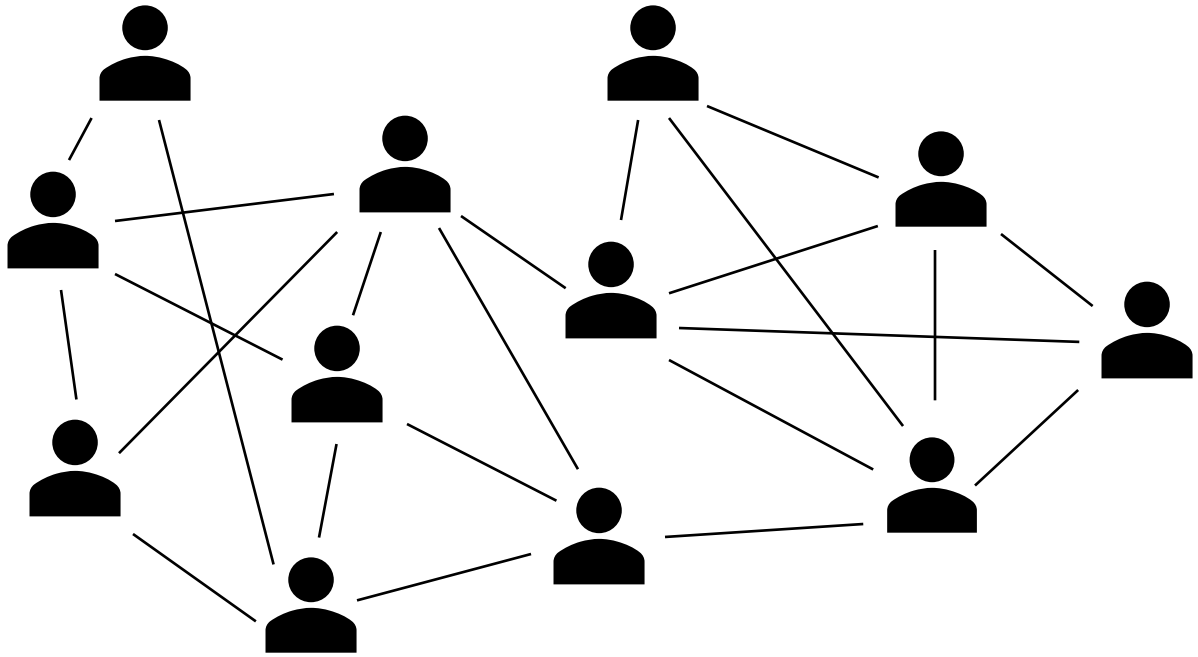

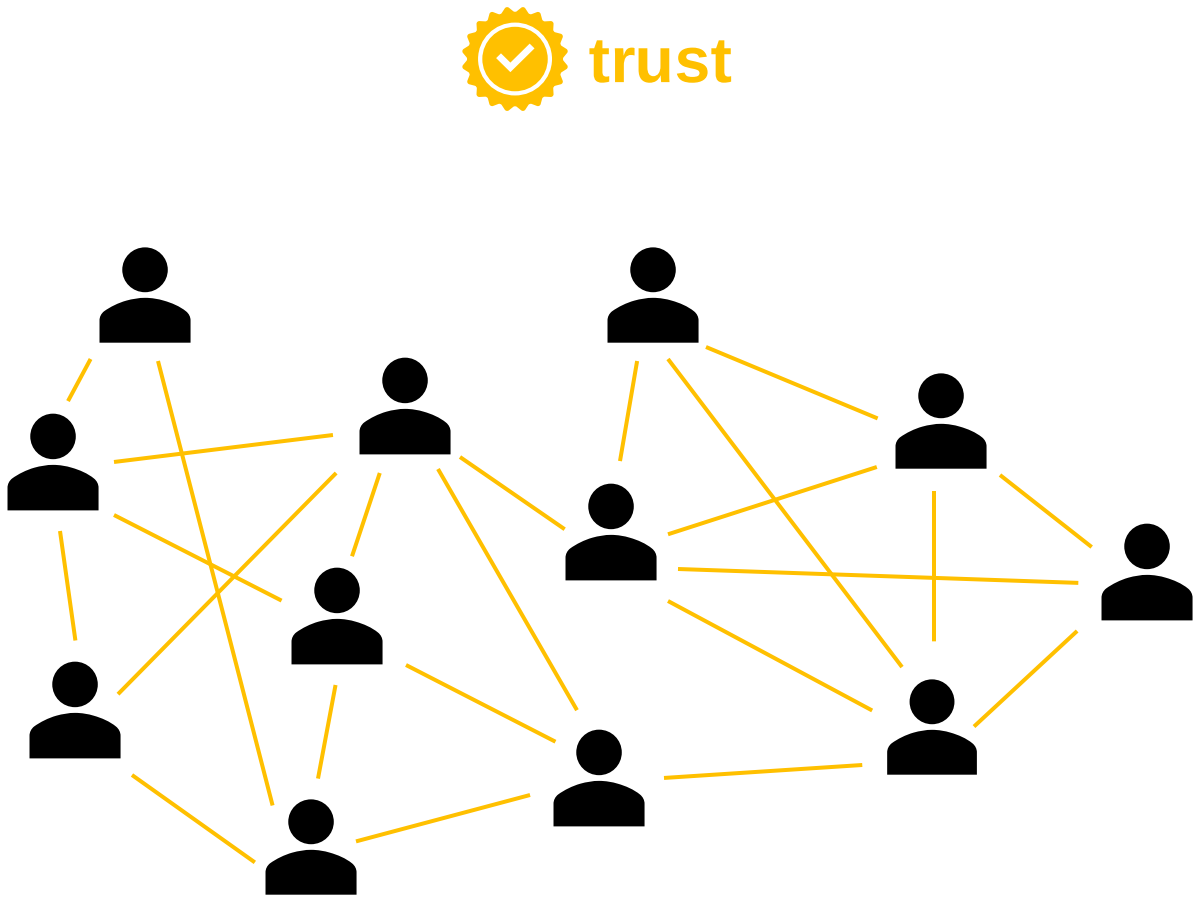
trust

Online Social Network (OSN)

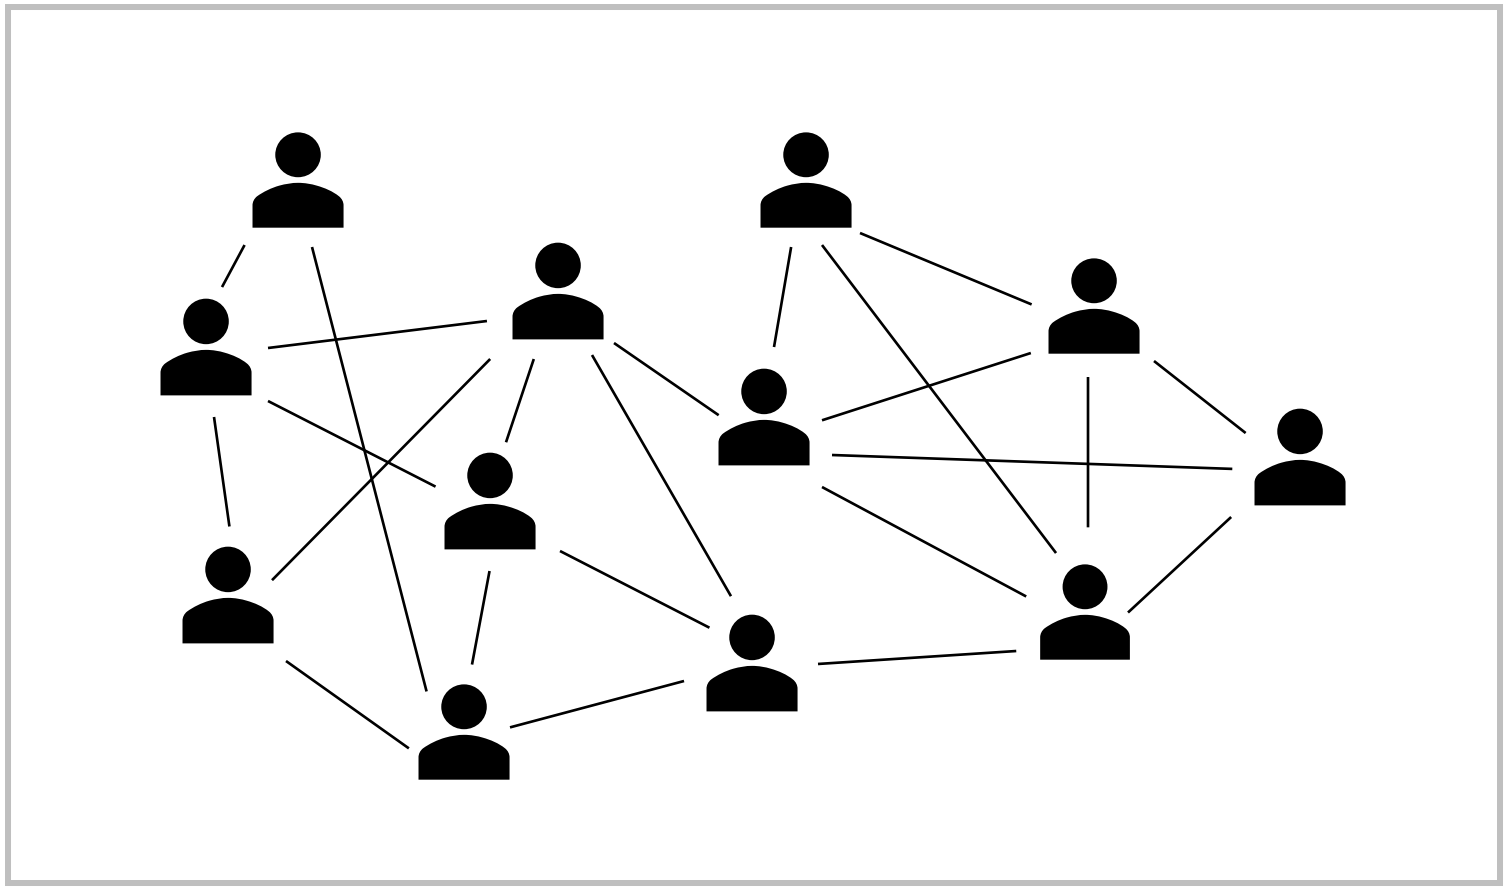

Users on a social media platform

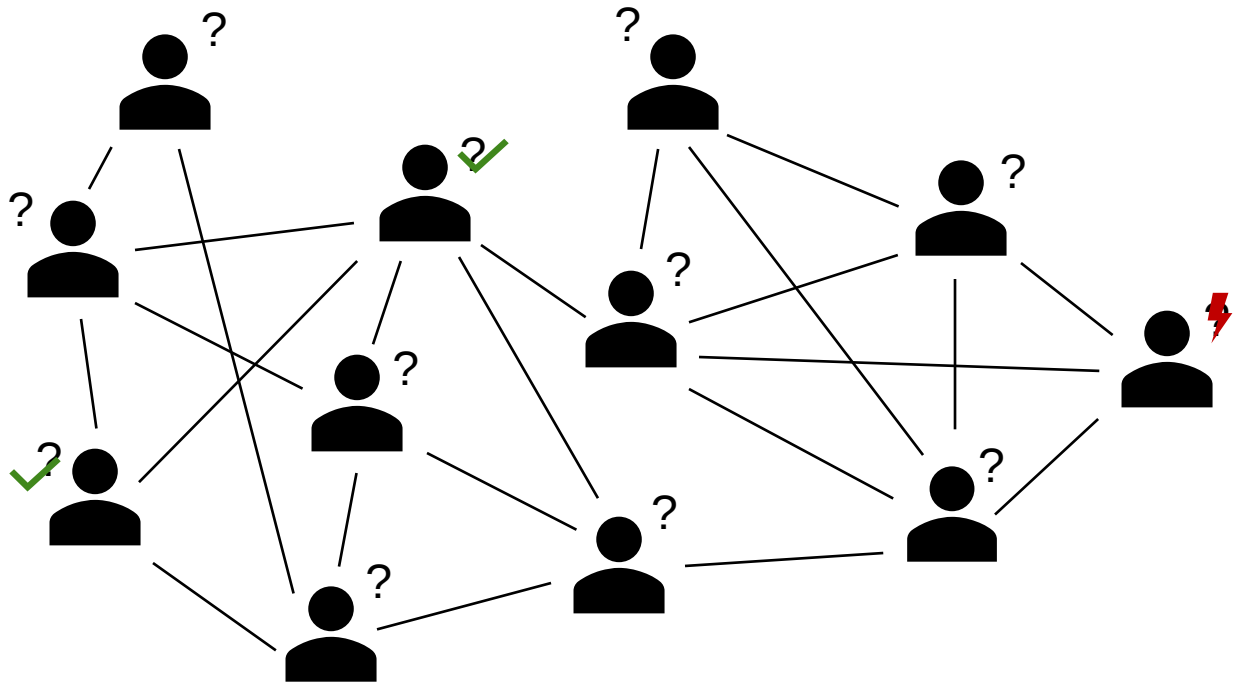

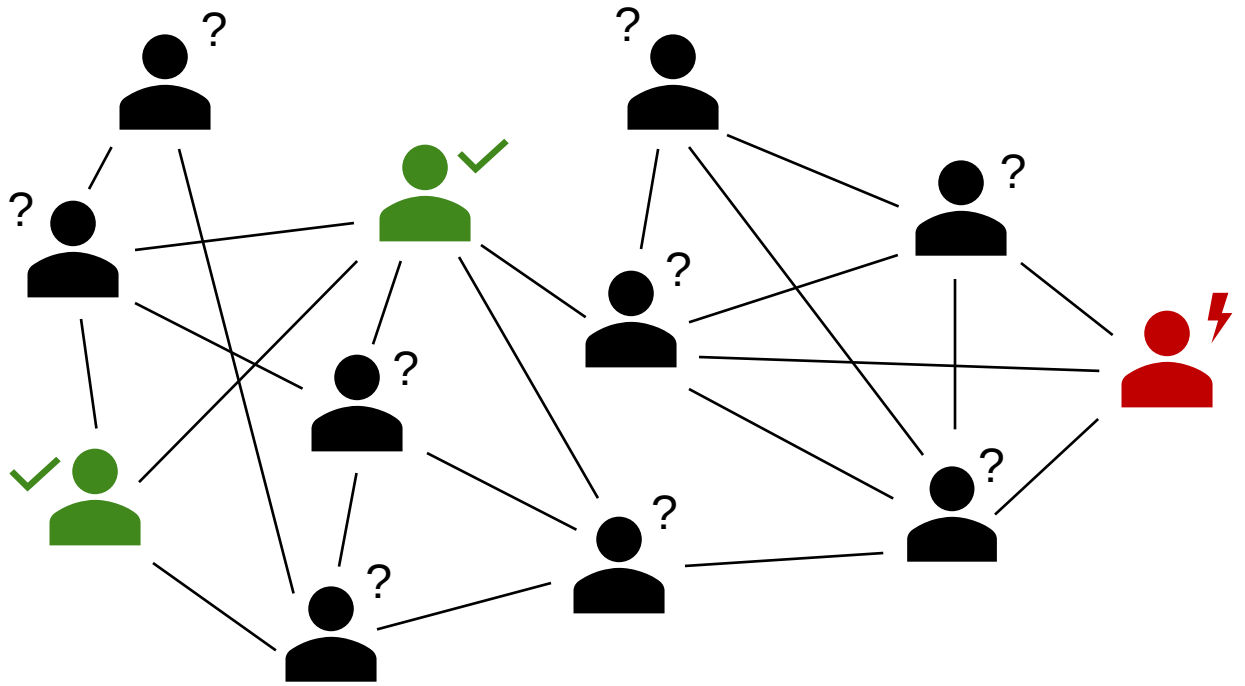

Background & Related Work

	SybilRank (2012)	**SybilBelief (2014)**	**SybilSCAR (2019)**
Approach	Random Walks (RW)	Loopy Belief Propagation (LBP)	Local rule-based propagation
Labels (honest / Sybil)	Only honest	Both	Both
Guaranteed convergence	Yes	No	Yes
Computational complexity	O(n log n)	O(n) per iteration	O(n log n)
Main advantage	Computationally efficient	Can use both label types	Combines benefits of RW and LBP

Modern Sybils

«The Eiffel Tower is a famous landmark in Paris, France. It stands tall at 330 meters (1,083 feet) with a square base. It is named after the engineer Gustav Eiffel.»

«The Eiffel Tower stands in Paris, France. It reaches a height of 330 meters (1,083 feet). The tower is named after Gustave Eiffel, the engineer whose company designed and built it.»

Claude 3.5 Sonnet
«write exactly 3 very short sentences about the eiffel tower, including its location, height and name origin»

Results: Pre-training on Small Network

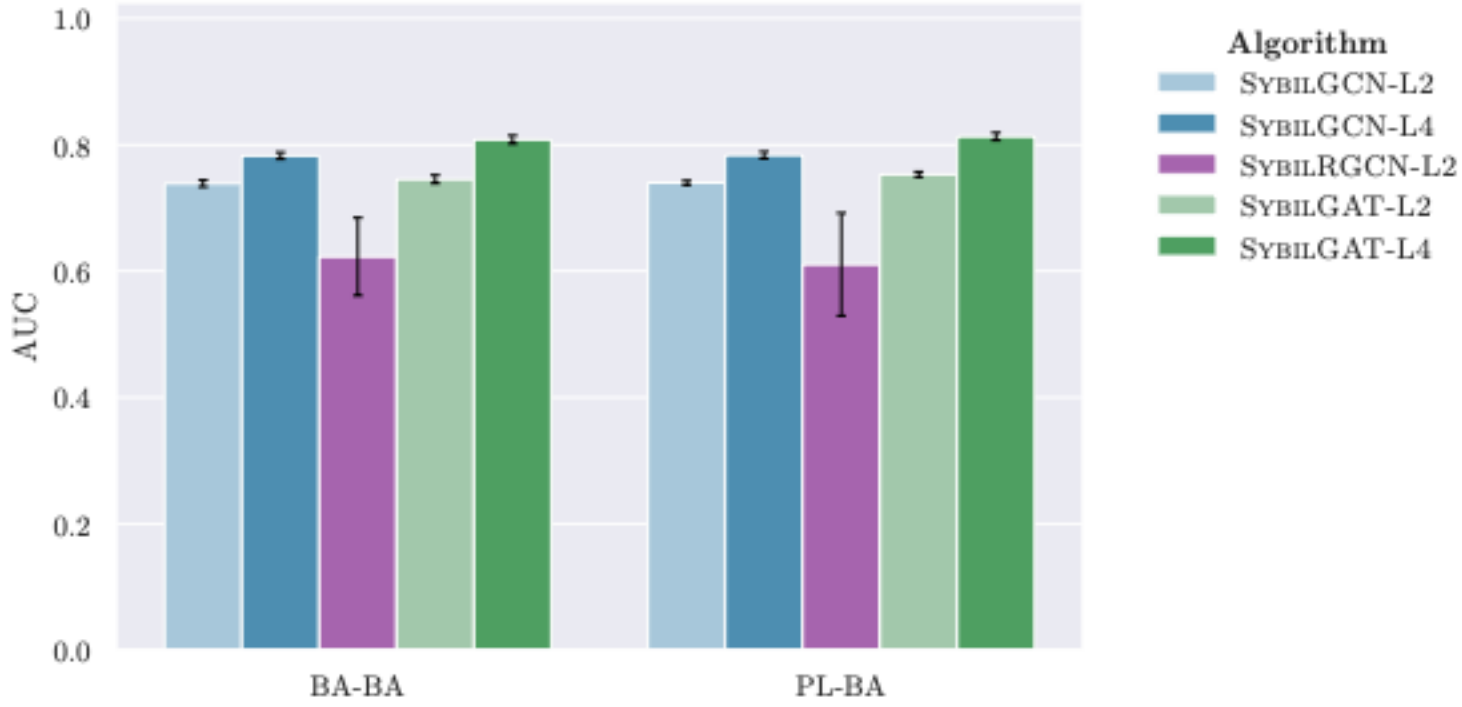

Results: Pre-training on Small Network

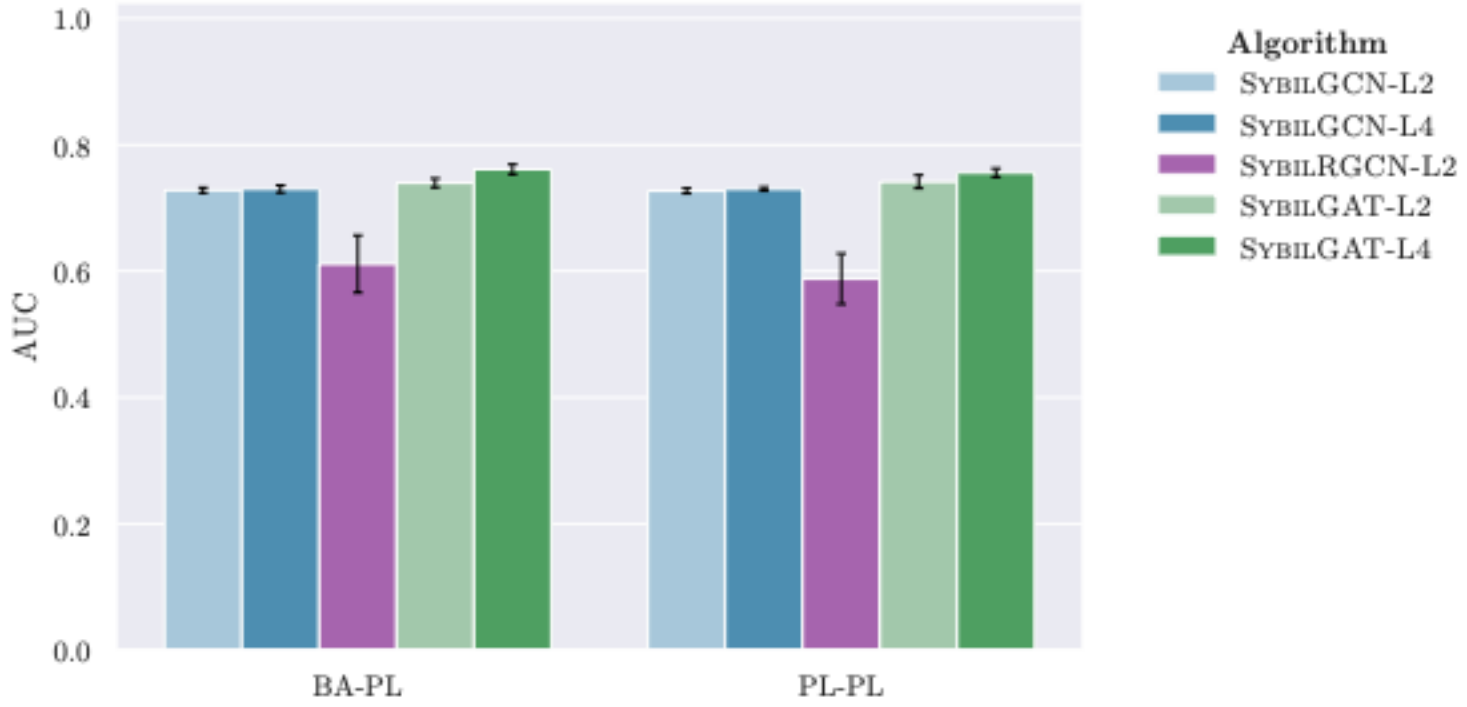

Results: Pre-training on Small Network

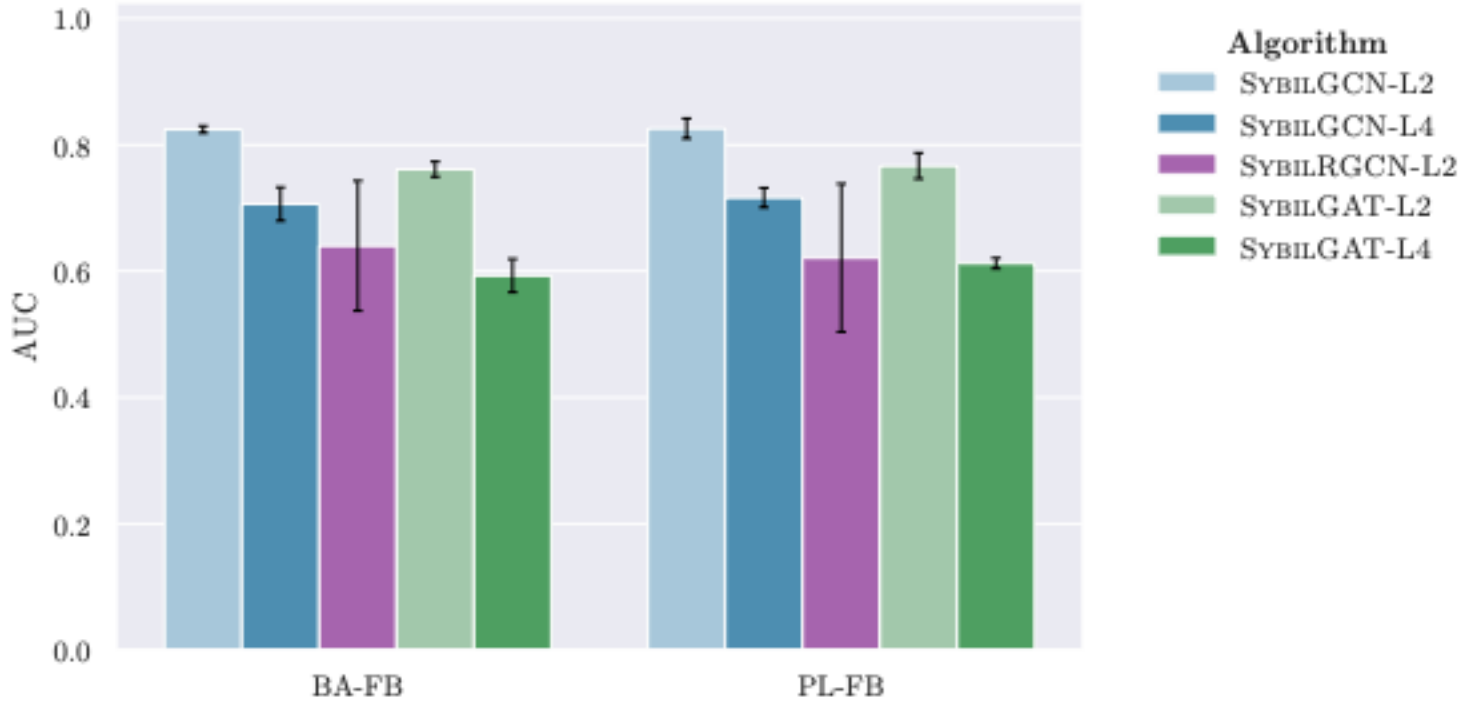

Results: Robustness
Data Model

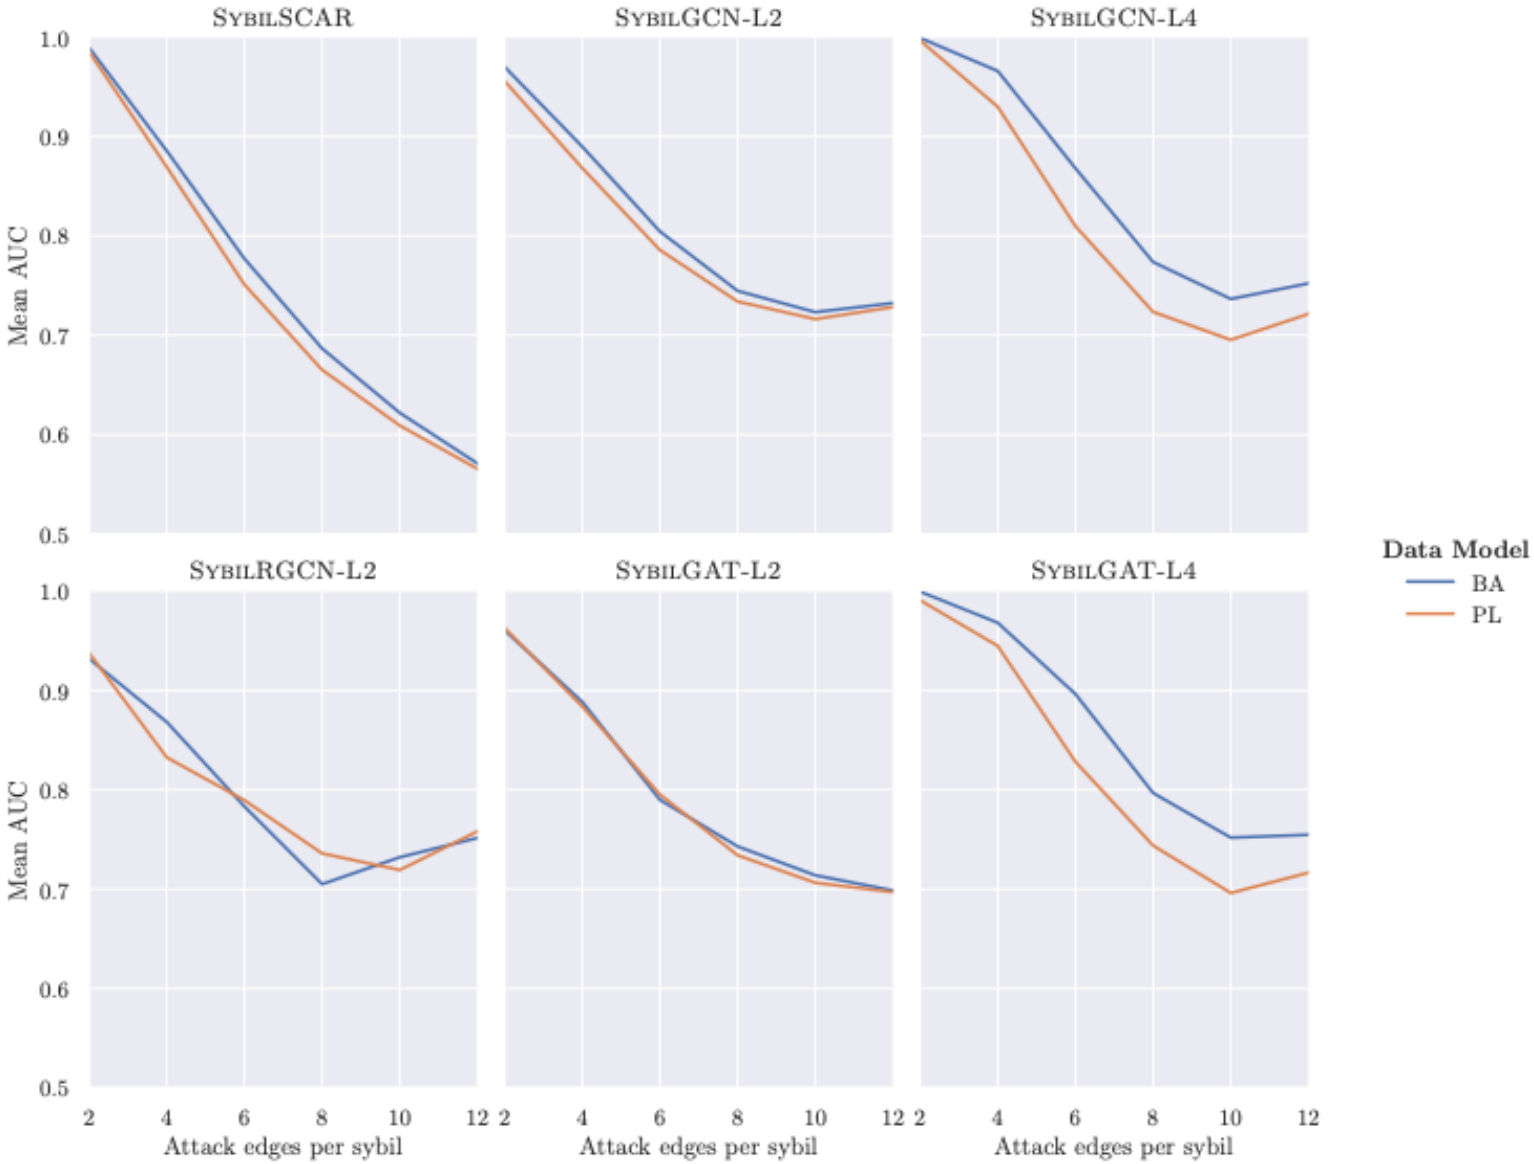

Results: Robustness
Data Model

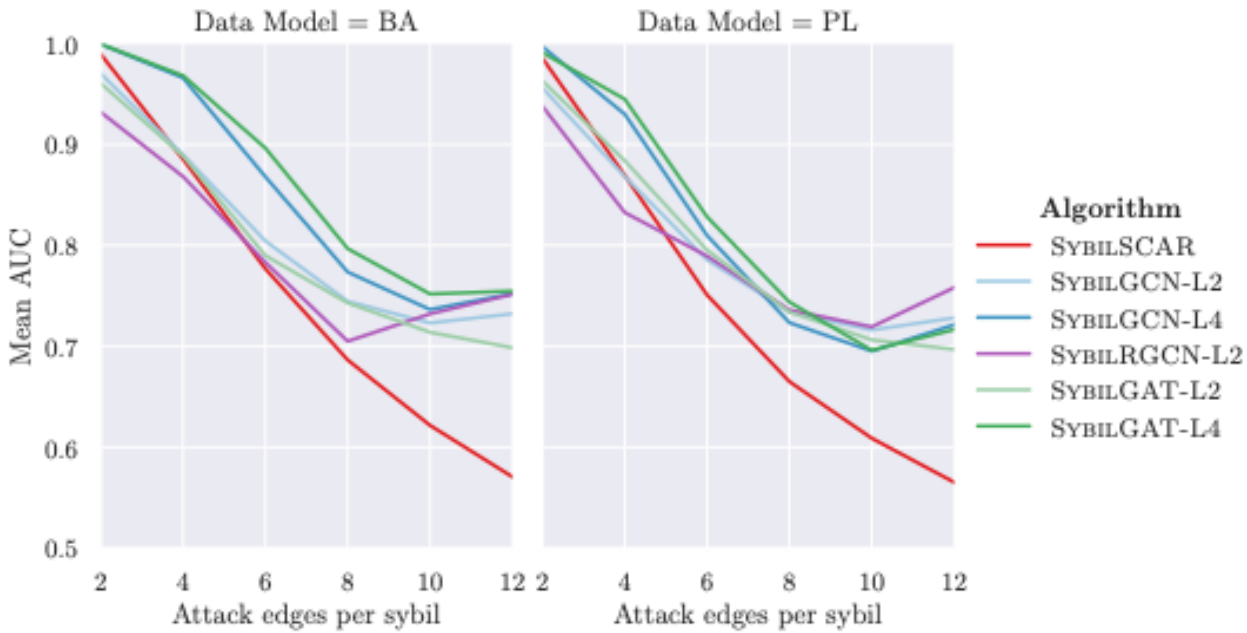

Results: Robustness

Network Size

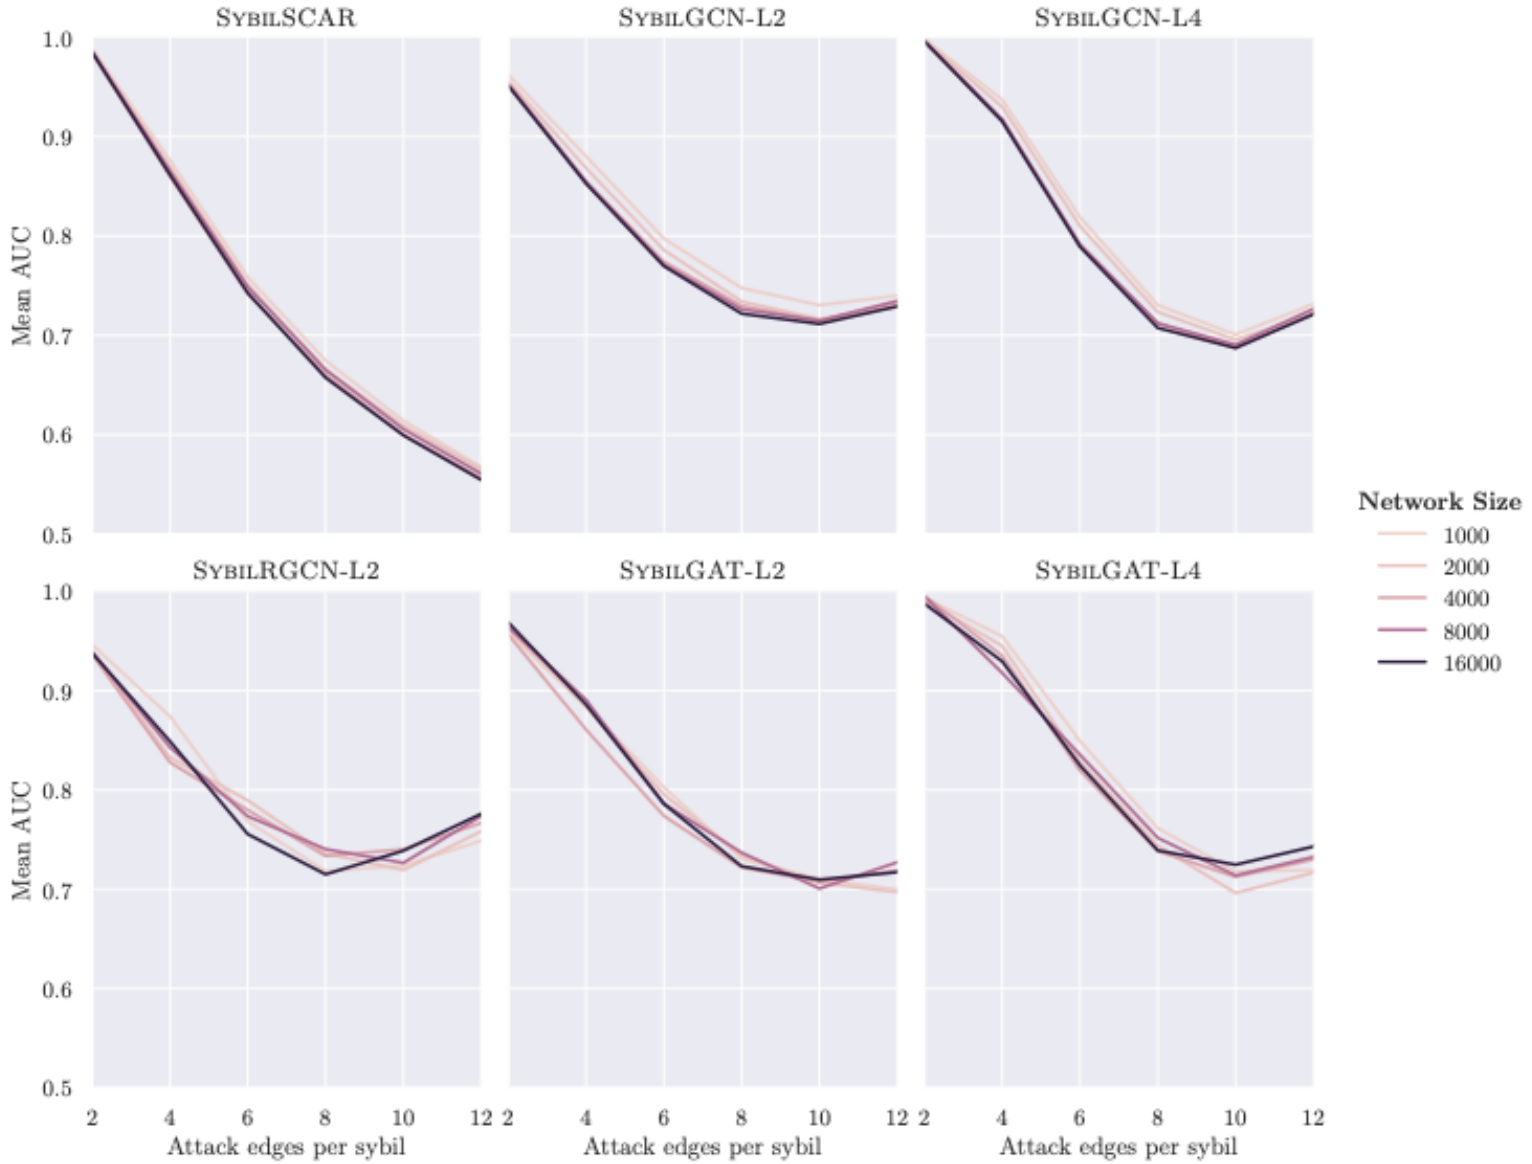

Results: Robustness

Network Size

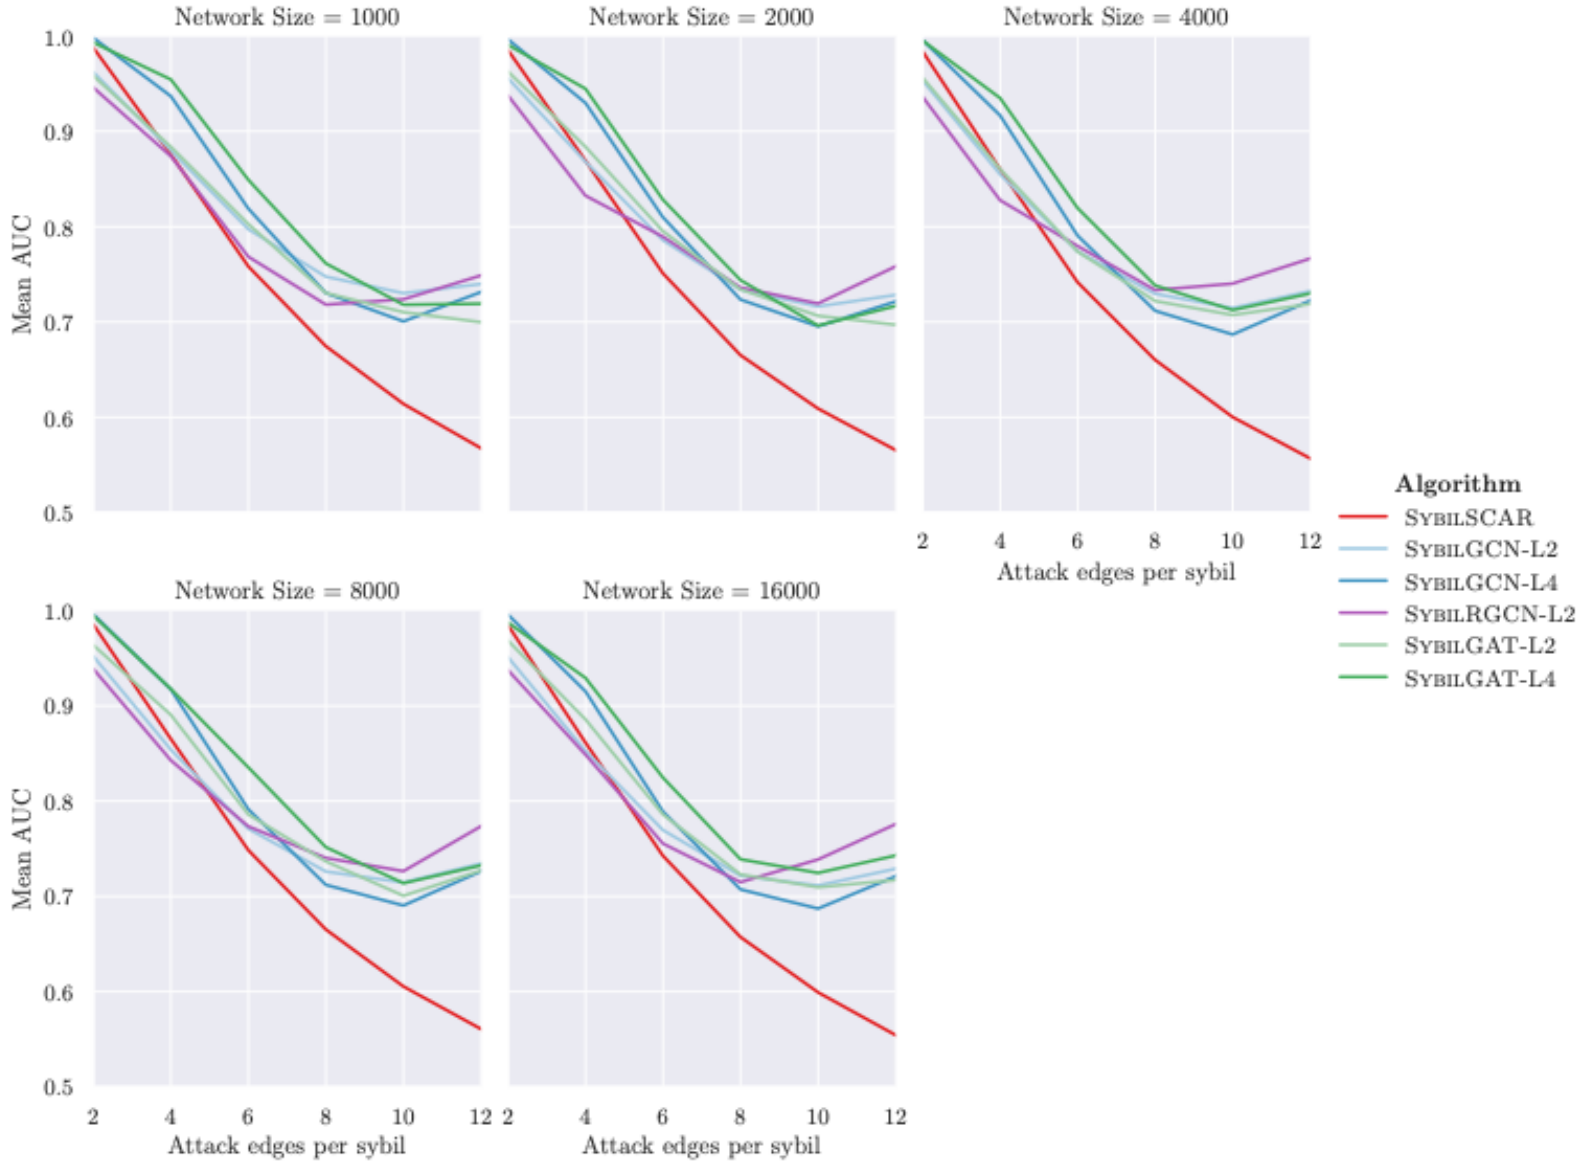

Results: Robustness
Training Set Size

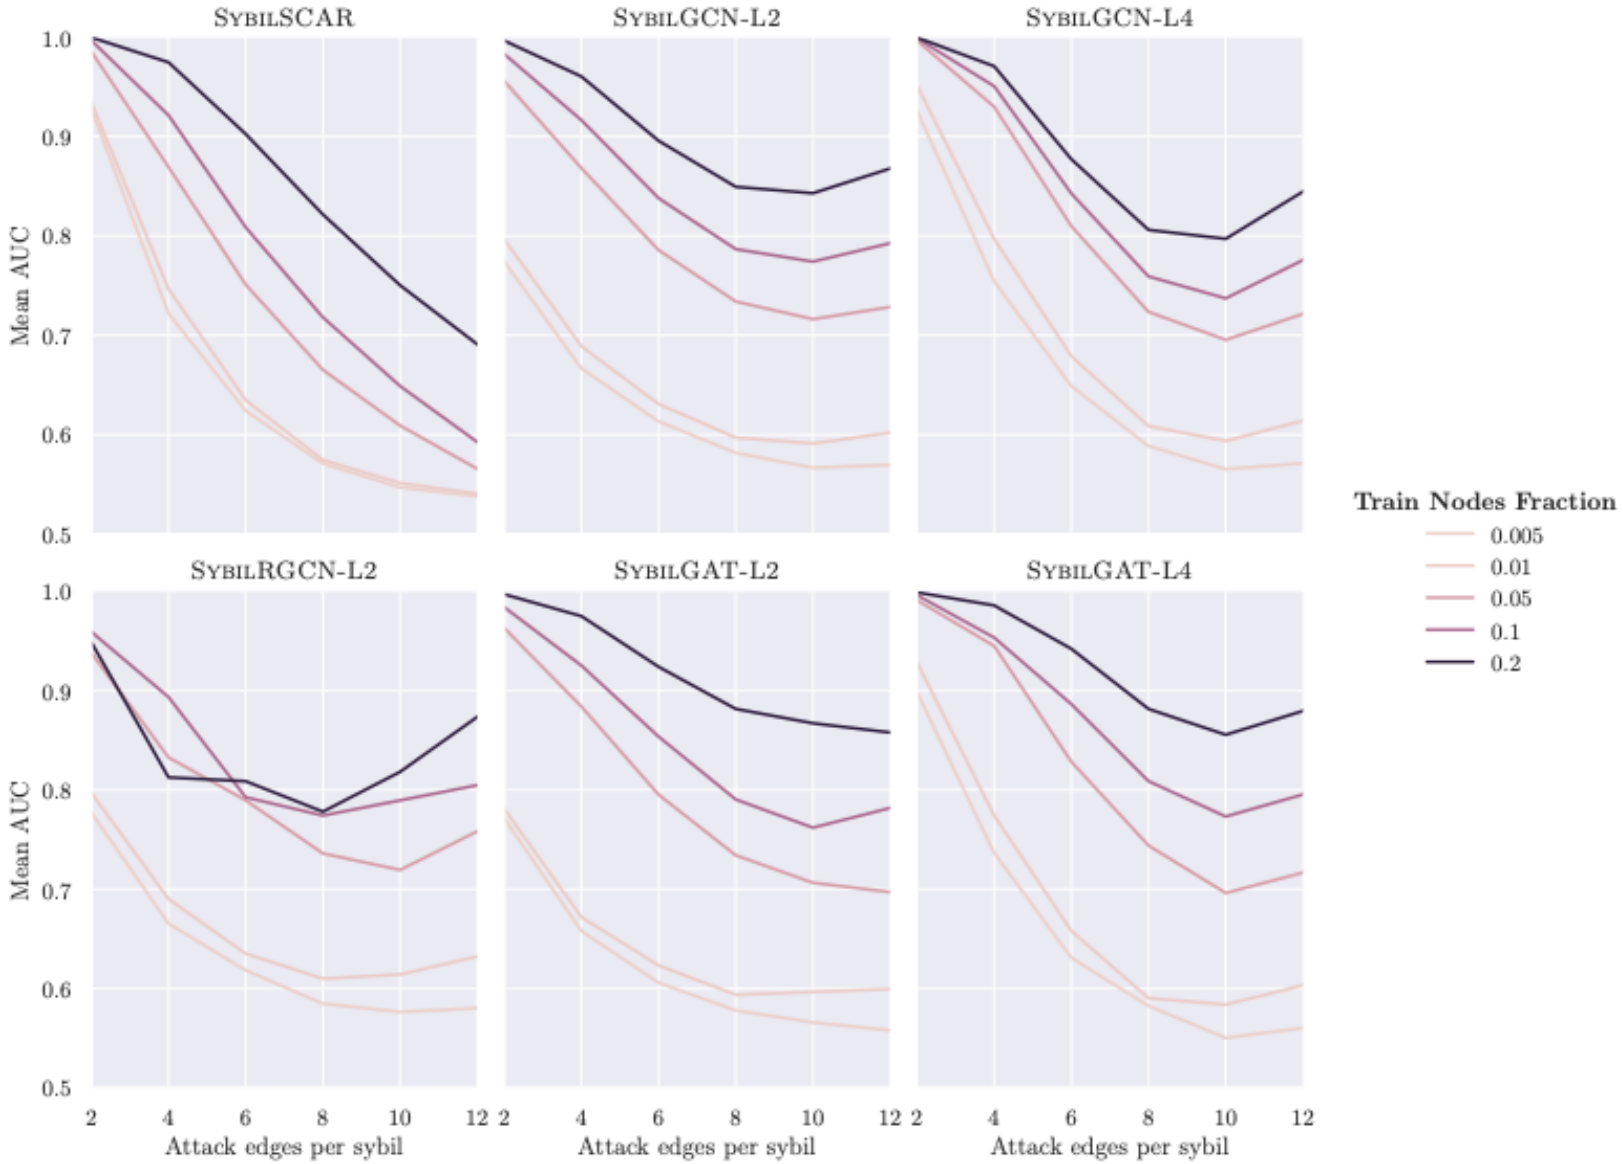

Results: Robustness
Training Set Size

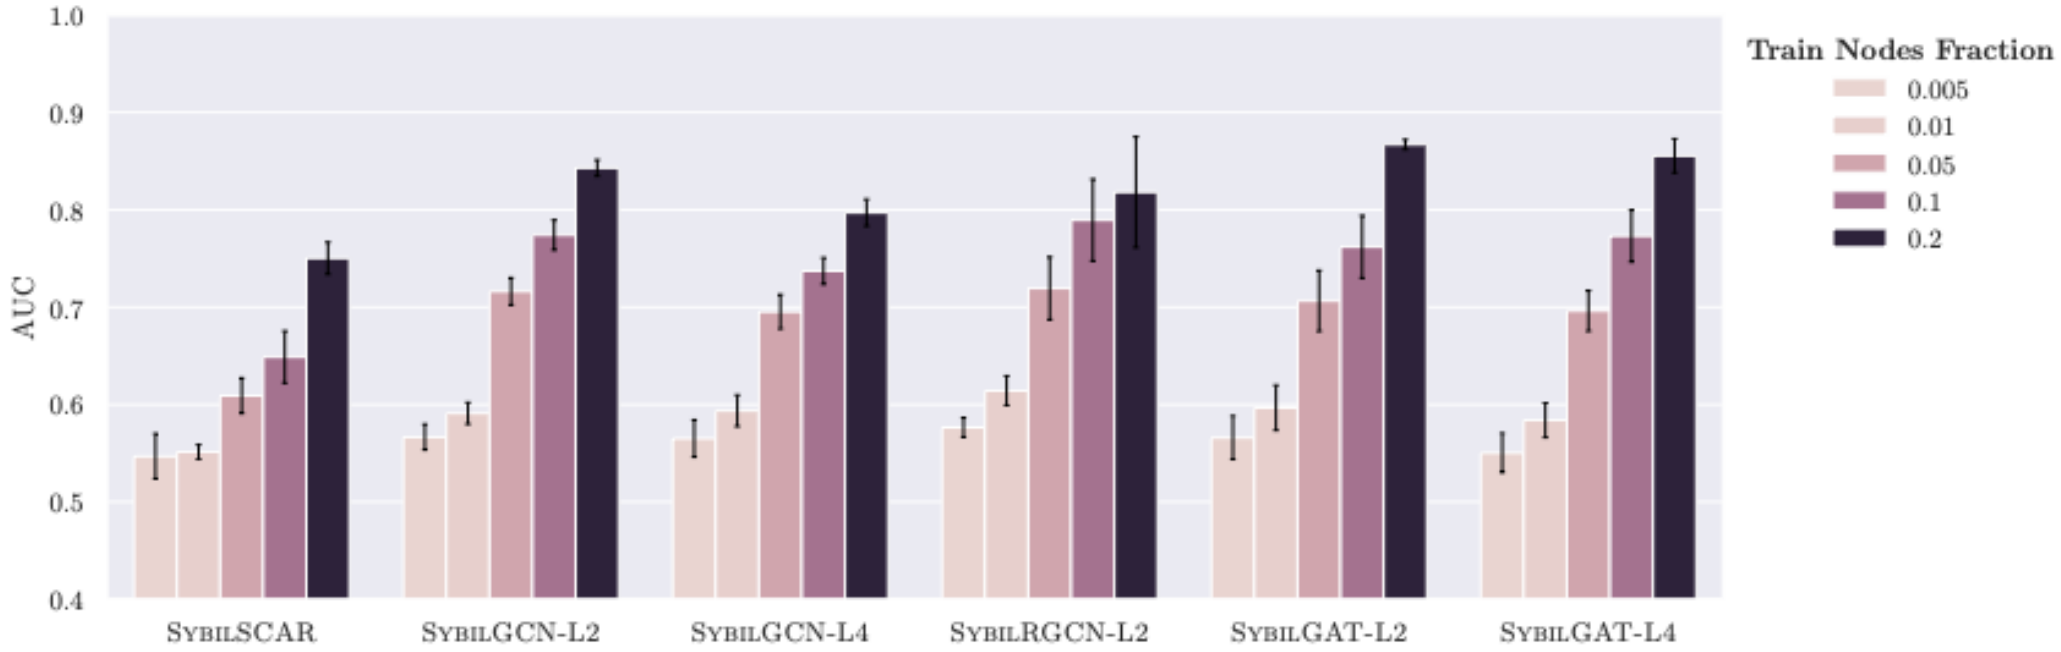

Results: Robustness

Training Set Size

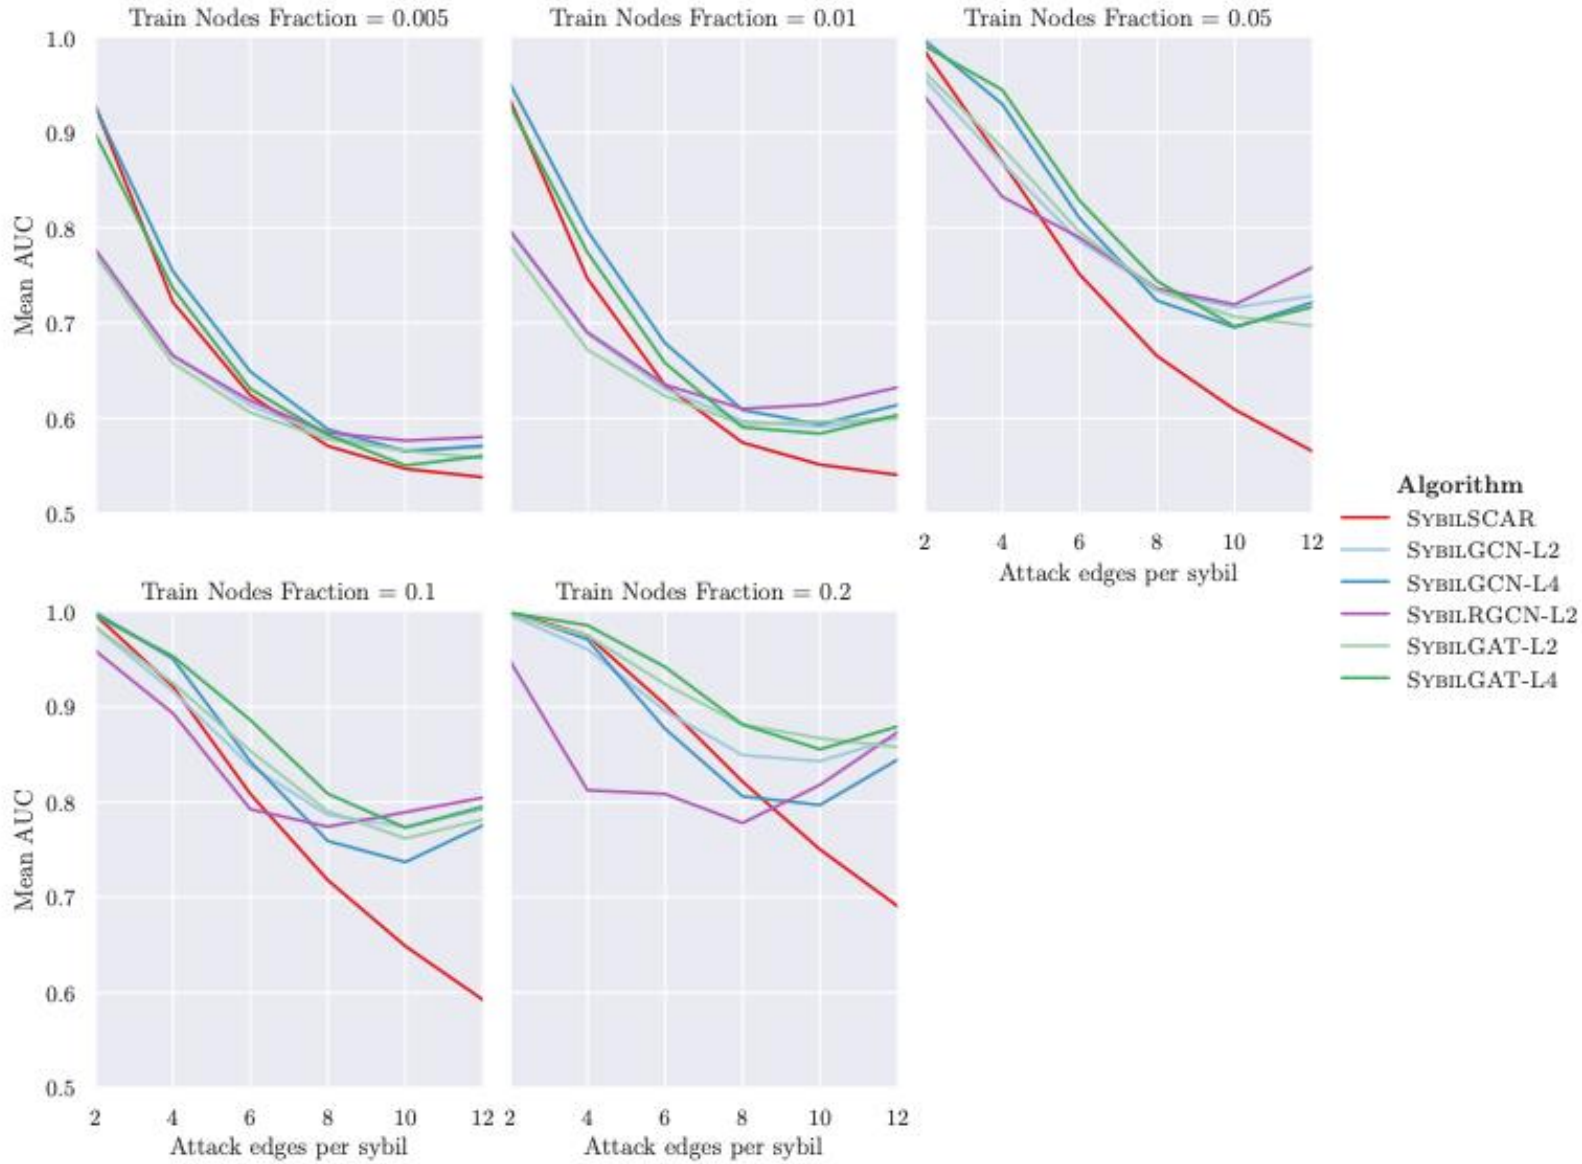

Results: Robustness

Label Noise Level

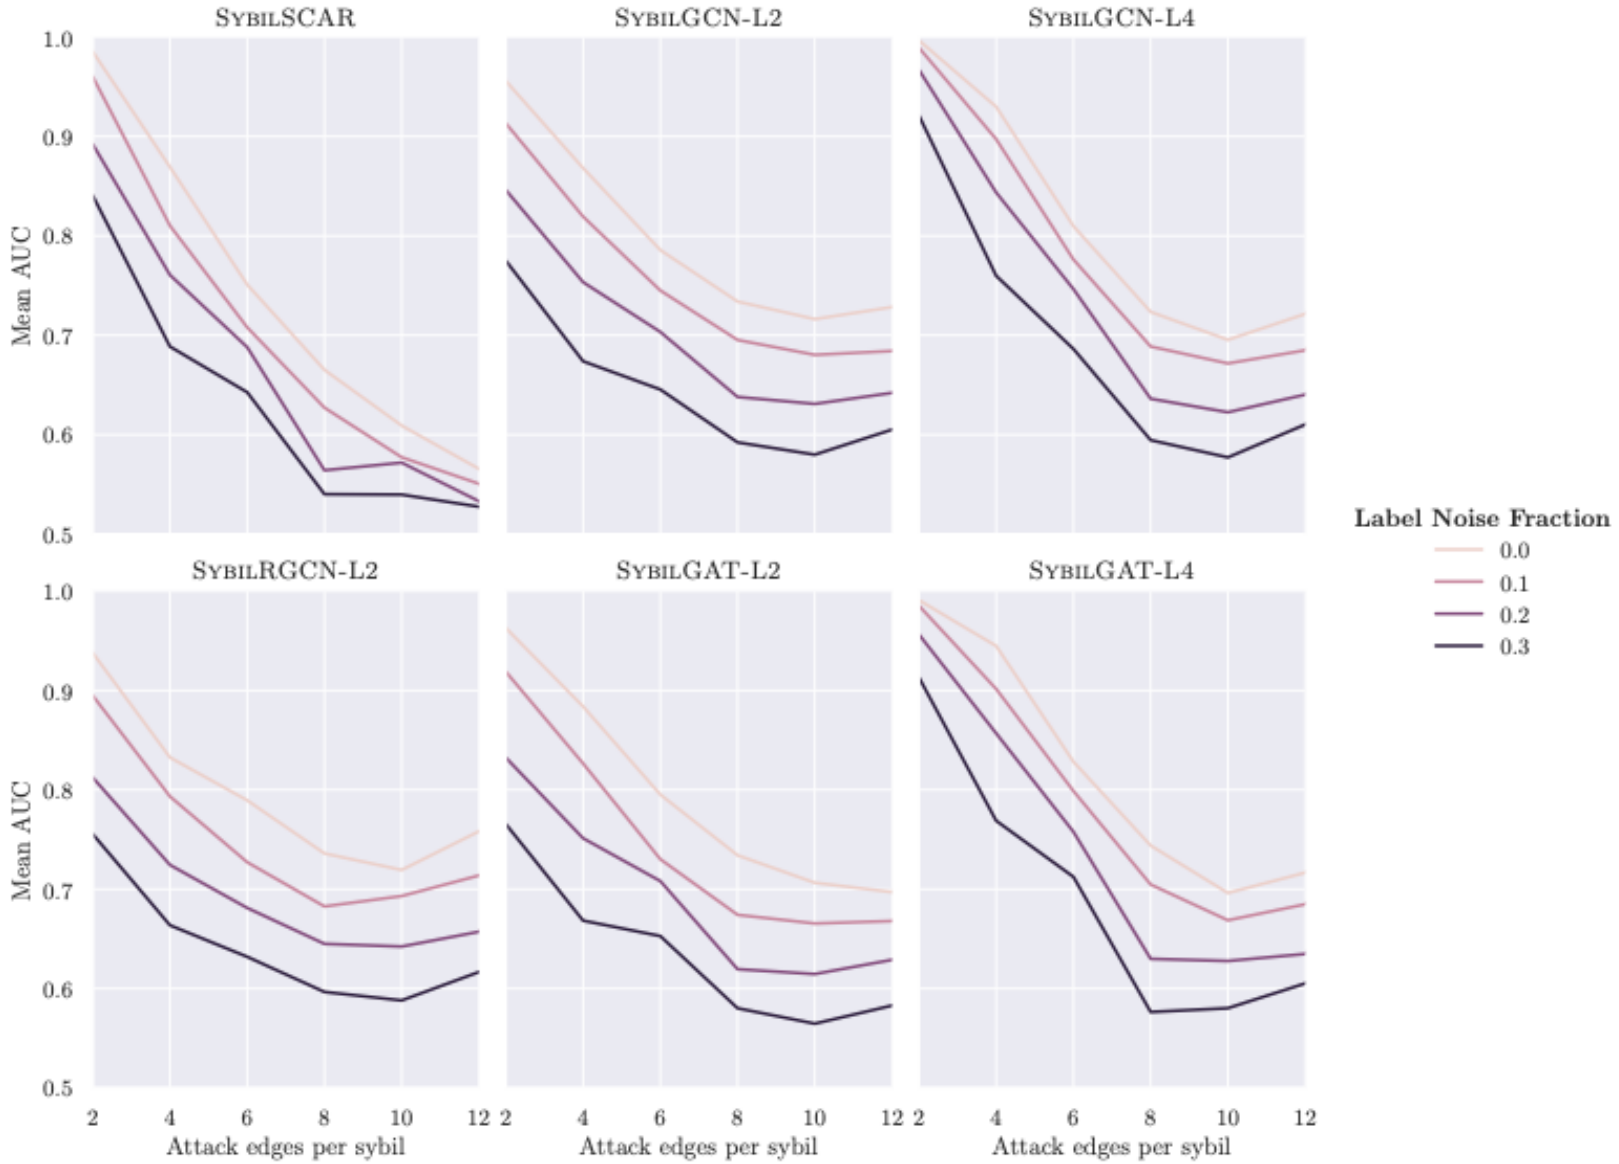

Adversarial Attack

Random attack edge

With prob. $1 - p_t$

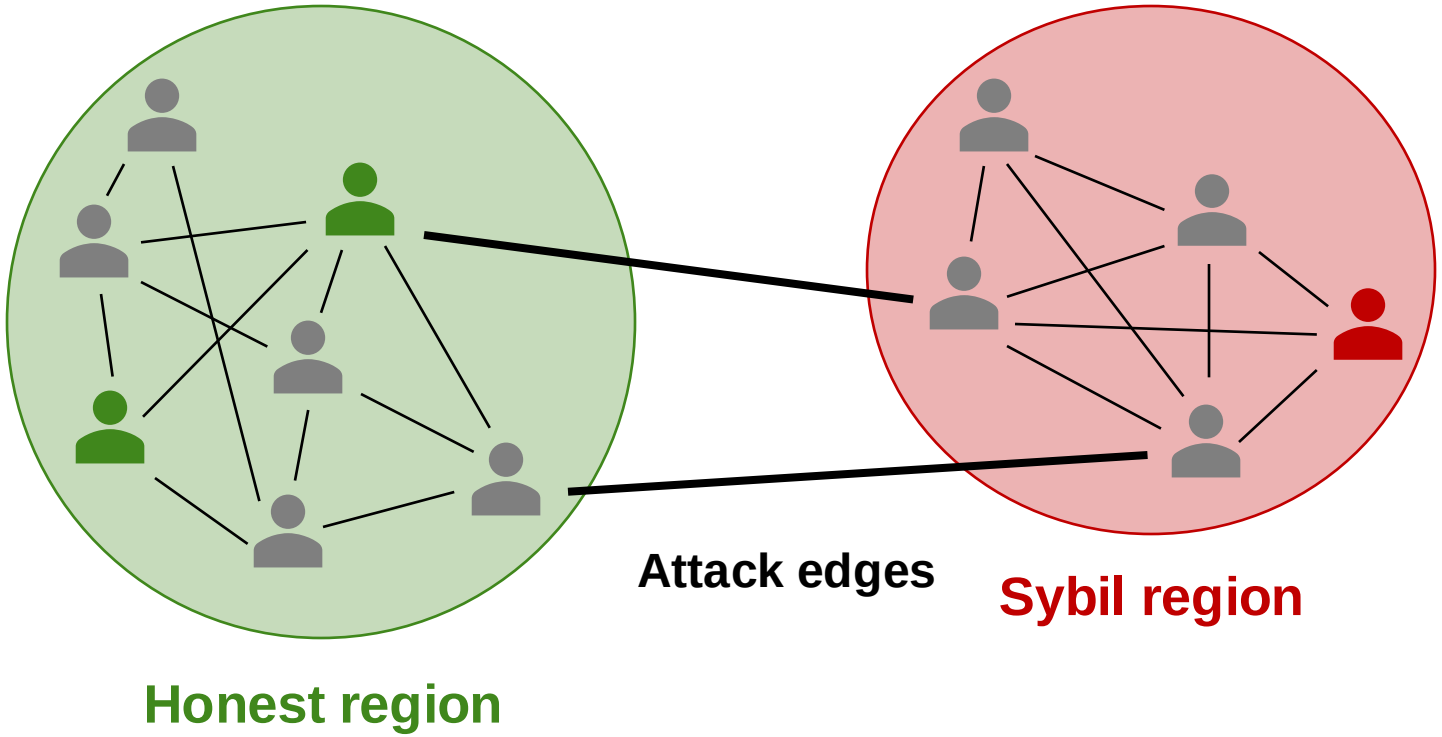

Adversarial Attack

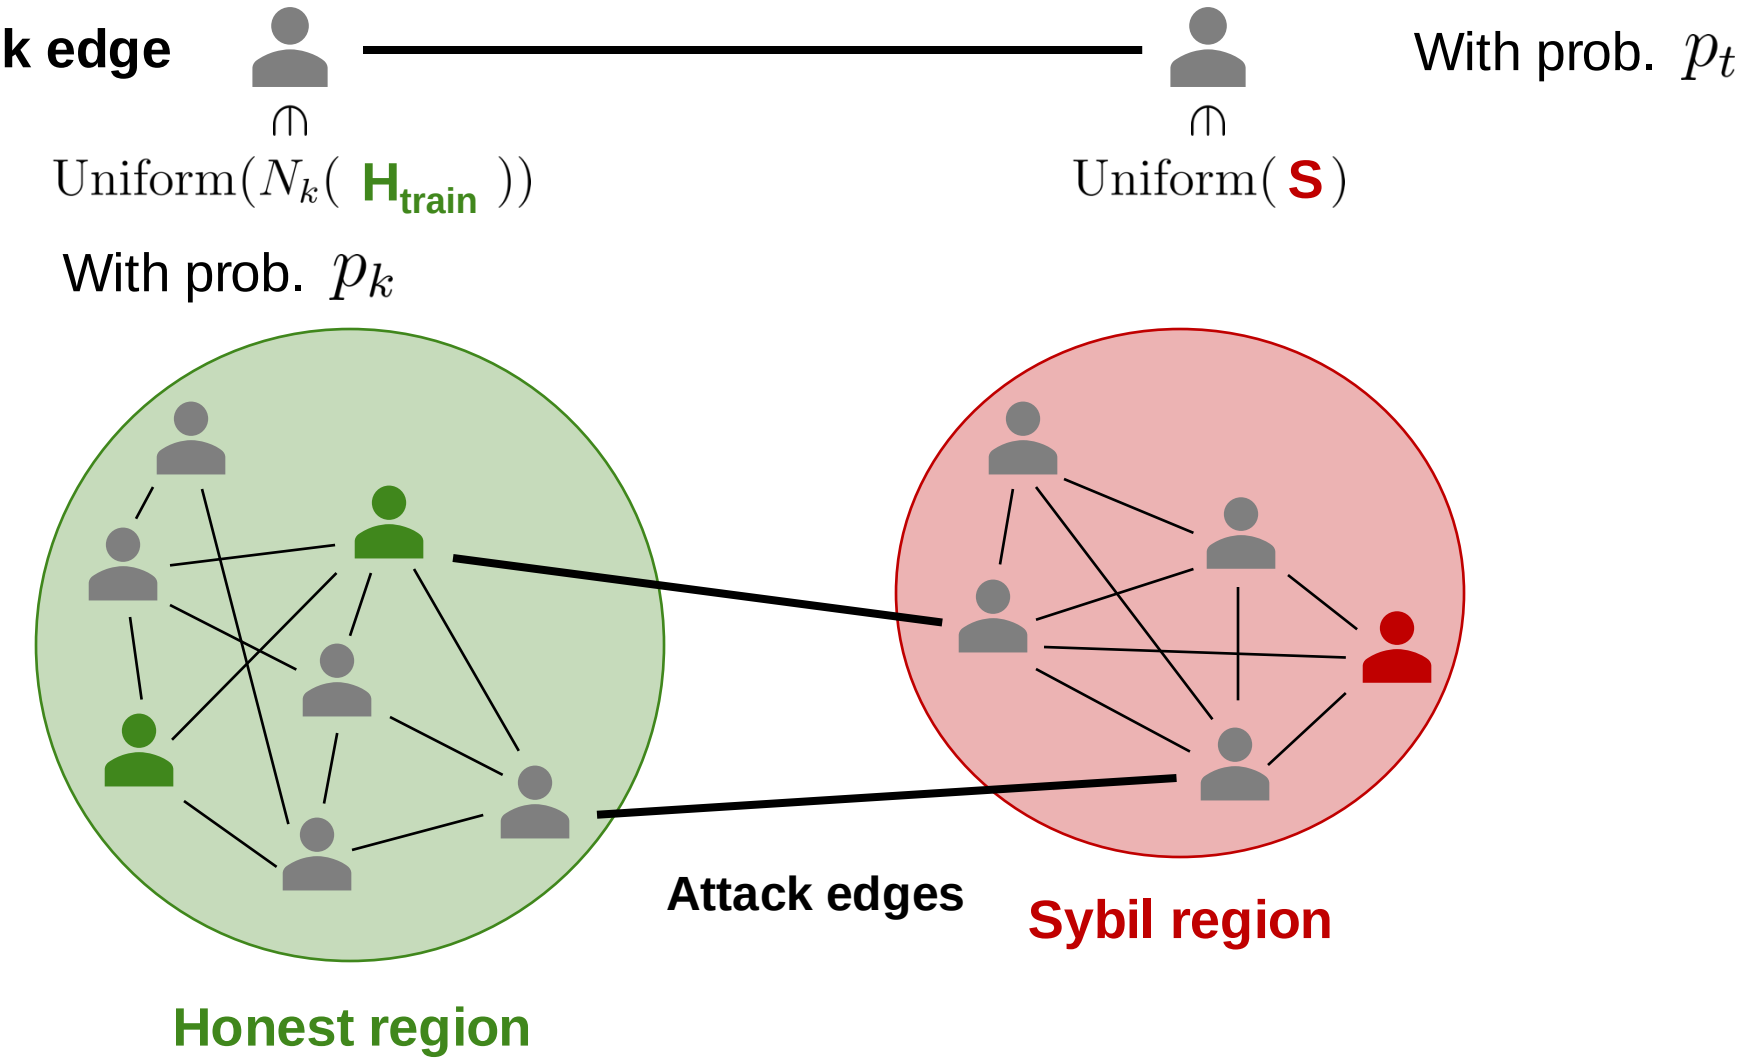

Results: Adversarial Attacks

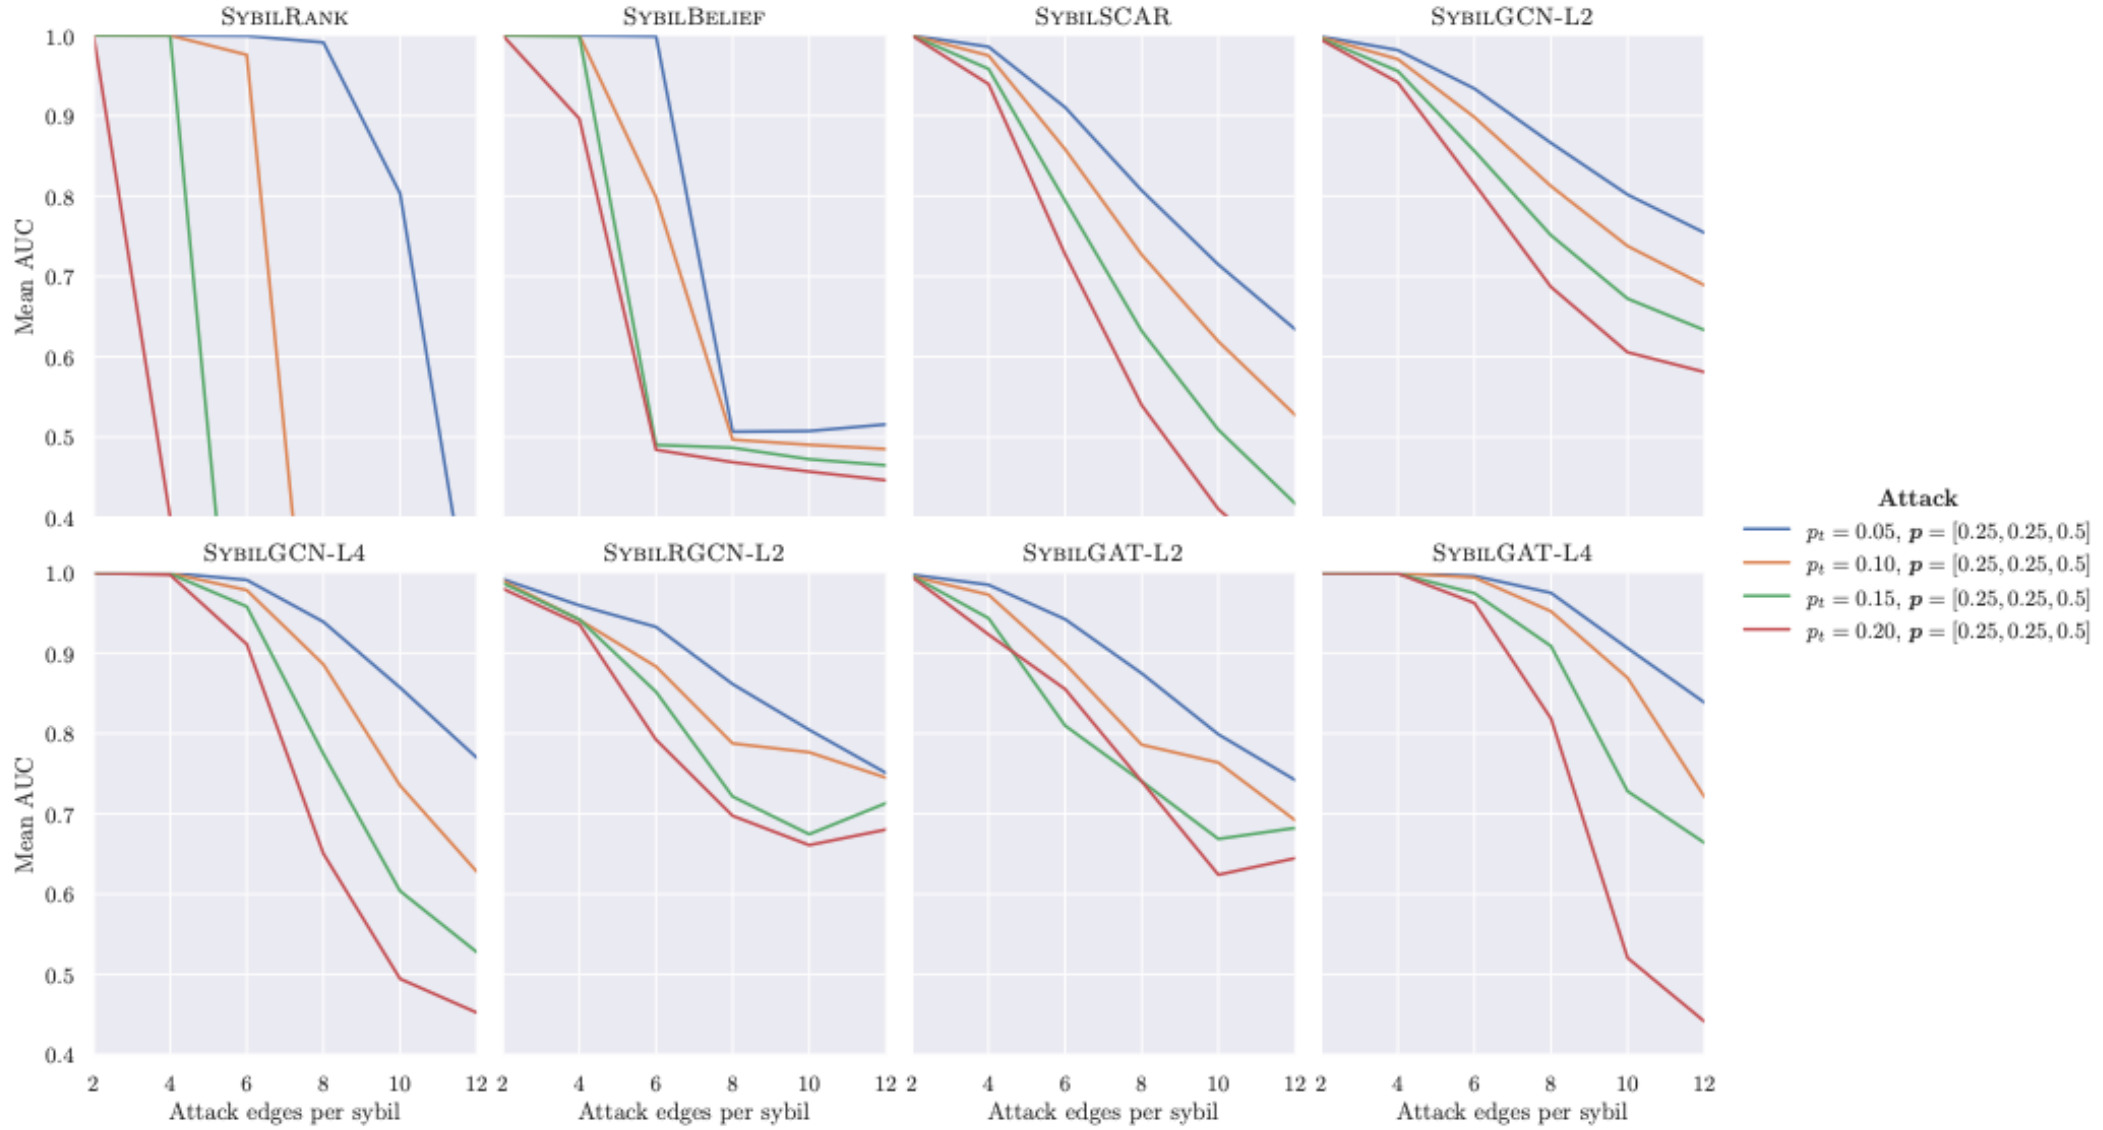

Results: Pre-training Before Attack